

CUNY SPS DATA 621 - CTG5 - HW1

Betsy Rosalen, Gabrielle Bartomeo, Jeremy O'Brien, Lidiia Tronina, Rose Koh

February 27, 2019

Contents

1	DATA EXPLORATION	1
1.1	Instructions: (Betsy: we can take these out later, but it helps me to see them right here for now...)	1
1.2	Possible writeup for Data Exploration	1
2	DATA PREPARATION	11
2.1	Instructions:	11
2.2	Possible writeup for Data Preparation	11
2.3	Information in general, don't use the writeups here as-is	12
2.4	Missing Values	12
2.5	NA Imputation	12
2.6	Feature Engineering	13
3	BUILD MODELS	14
3.1	Instructions:	14
3.2	MODEL 1	14
3.3	MODEL 2	19
3.4	MODEL 3	31
4	SELECT MODELS	33
4.1	Instructions:	33

1 DATA EXPLORATION

1.1 Instructions: (Betsy: we can take these out later, but it helps me to see them right here for now...)

Describe the size and the variables in the moneyball training data set. Consider that too much detail will cause a manager to lose interest while too little detail will make the manager consider that you aren't doing your job. Some suggestions are given below. Please do NOT treat this as a check list of things to do to complete the assignment. You should have your own thoughts on what to tell the boss. These are just ideas.

- a. Mean / Standard Deviation / Median
- b. Bar Chart or Box Plot of the data and/or Histograms
- c. Is the data correlated to the target variable (or to other variables?)
- d. Are any of the variables missing and need to be imputed "fixed"?

1.2 Possible writeup for Data Exploration

Professionals and gamblers alike are always seeking to optimize their chances of winning, whether it be sports, games, or their bets on them. Major League Baseball is a multibillion dollar industry where individual teams, players, and those who profit off of their success stand to benefit most from such optimization.

In order to determine the best way to infer whether the 162 games in a baseball team's year will result in more wins overall, data from 1871 to 2006 where each set of values represented a season for an unnamed team, totalling 2,276 records. For each team their number of wins in a given year were given with a maximum possible of 162 wins, in addition to that team's base hits, doubles, triples, homeruns, walks, and strikeouts by batters, batters hit by pitches, bases stolen by batters and the number of times they were caught stealing, the number of errors, double plays, walks, hits, and homeruns allowed, and strikeouts by pitchers.

1.2.1 Data before imputing values

	n	min	mean	median	max	sd
TARGET_WINS	2276	0	80.79086	82.0	146	15.75215
BATTING_H	2276	891	1469.26977	1454.0	2554	144.59120
BATTING_2B	2276	69	241.24692	238.0	458	46.80141
BATTING_3B	2276	0	55.25000	47.0	223	27.93856
BATTING_HR	2276	0	99.61204	102.0	264	60.54687
BATTING_BB	2276	0	501.55888	512.0	878	122.67086
BATTING_SO	2174	0	735.60534	750.0	1399	248.52642
BASERUN_SB	2145	0	124.76177	101.0	697	87.79117
BASERUN_CS	1504	0	52.80386	49.0	201	22.95634
PITCHING_H	2276	1137	1779.21046	1518.0	30132	1406.84293
PITCHING_HR	2276	0	105.69859	107.0	343	61.29875
PITCHING_BB	2276	0	553.00791	536.5	3645	166.35736
PITCHING_SO	2174	0	817.73045	813.5	19278	553.08503
FIELDING_E	2276	65	246.48067	159.0	1898	227.77097
FIELDING_DP	1990	52	146.38794	149.0	228	26.22639

1.2.2 Data after imputing values

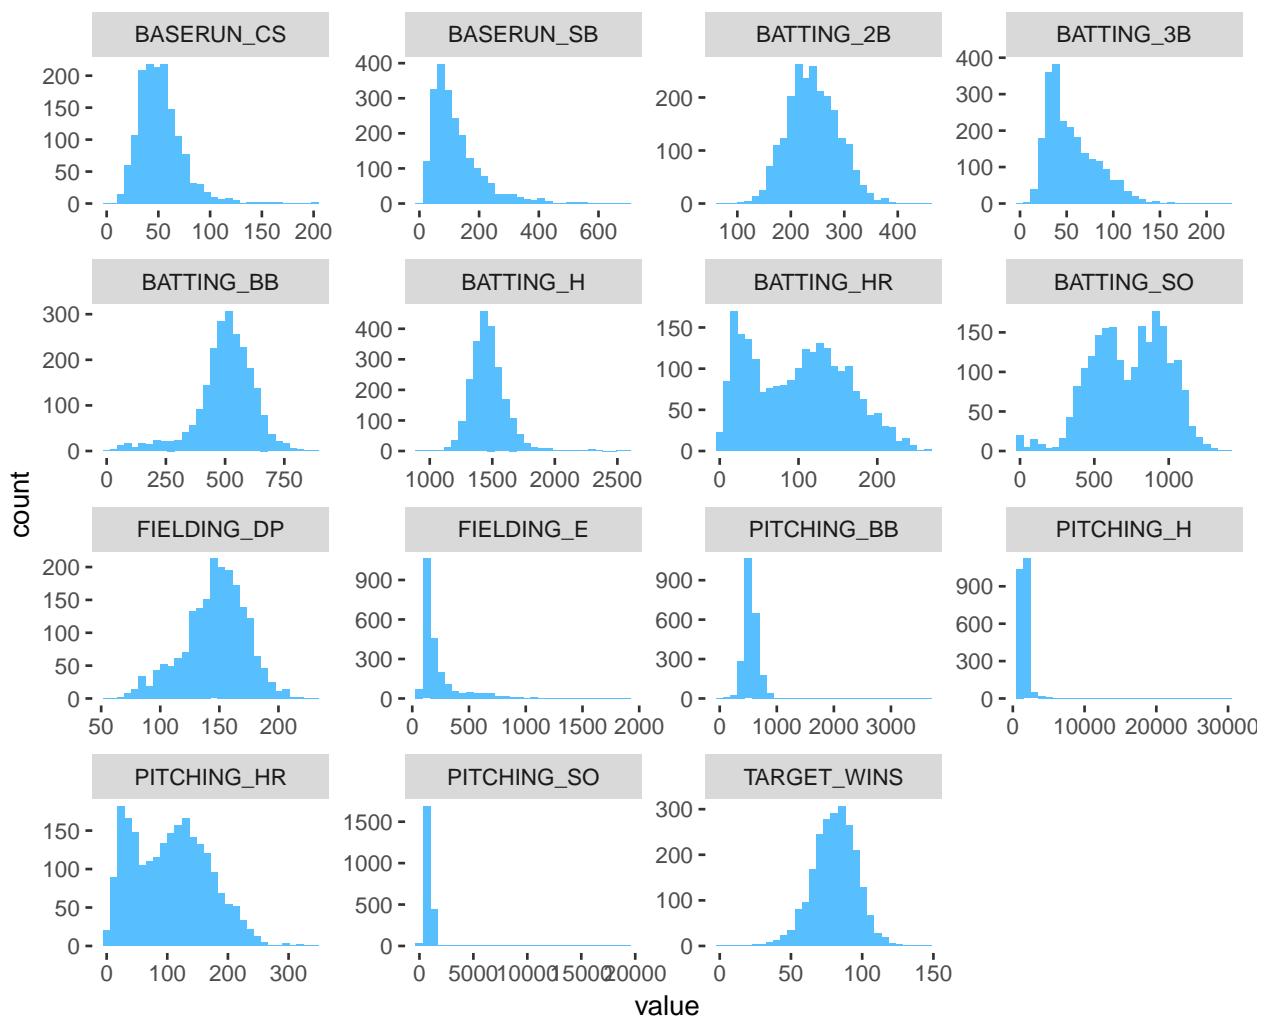
	n	min	mean	median	max	sd
TARGET_WINS	2276	0	80.79086	82.00000	146	15.75215
BATTING_H	2276	891	1469.26977	1454.00000	2554	144.59120
BATTING_2B	2276	69	241.24692	238.00000	458	46.80141
BATTING_3B	2276	0	55.25000	47.00000	223	27.93856
BATTING_HR	2276	0	99.61204	102.00000	264	60.54687
BATTING_BB	2276	0	501.55888	512.00000	878	122.67086
BATTING_SO	2276	0	728.22986	739.00000	1399	246.65040
BASERUN_SB	2276	0	124.62700	105.00000	697	85.28361
BASERUN_CS	2276	0	70.73189	56.42774	201	38.01002
PITCHING_H	2276	1137	1779.21046	1518.00000	30132	1406.84293
PITCHING_HR	2276	0	105.69859	107.00000	343	61.29875
PITCHING_BB	2276	0	553.00791	536.50000	3645	166.35736
PITCHING_SO	2276	0	807.62995	803.50000	19278	543.08096
FIELDING_E	2276	65	246.48067	159.00000	1898	227.77097
FIELDING_DP	2276	52	145.30987	146.00000	228	24.89909

1.2.3 Difference between original and imputed data

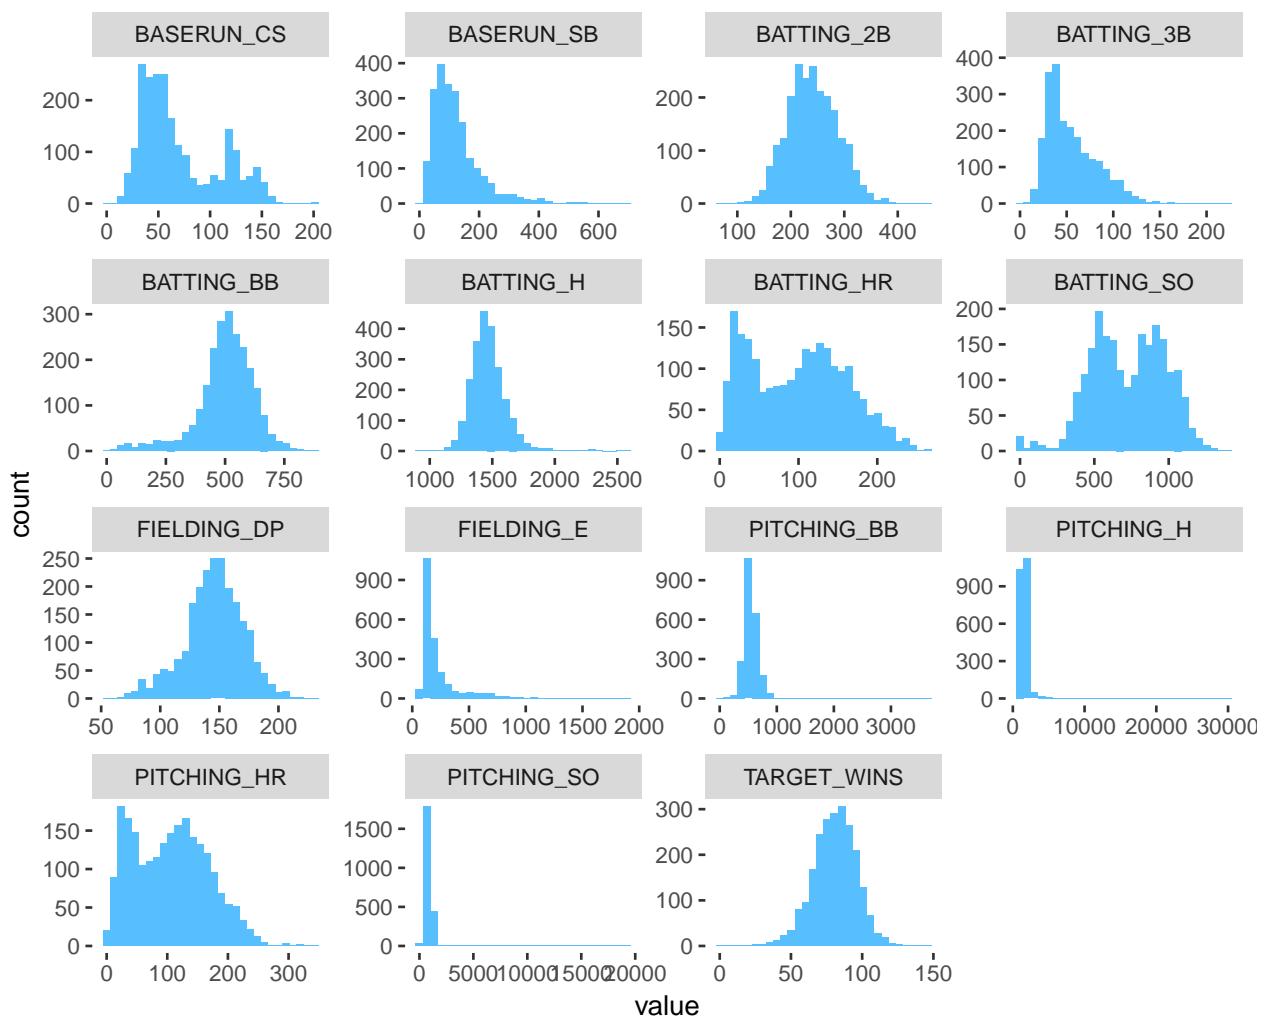
	n	min	mean	median	max	sd
TARGET_WINS	0	0	0.0000000	0.000000	0	0.000000
BATTING_H	0	0	0.0000000	0.000000	0	0.000000
BATTING_2B	0	0	0.0000000	0.000000	0	0.000000
BATTING_3B	0	0	0.0000000	0.000000	0	0.000000
BATTING_HR	0	0	0.0000000	0.000000	0	0.000000
BATTING_BB	0	0	0.0000000	0.000000	0	0.000000
BATTING_SO	-102	0	7.3754714	11.000000	0	1.876015
BASERUN_SB	-131	0	0.1347715	-4.000000	0	2.507556
BASERUN_CS	-772	0	-17.9280286	-7.427742	0	-15.053684
PITCHING_H	0	0	0.0000000	0.000000	0	0.000000
PITCHING_HR	0	0	0.0000000	0.000000	0	0.000000
PITCHING_BB	0	0	0.0000000	0.000000	0	0.000000
PITCHING_SO	-102	0	10.1004968	10.000000	0	10.004072
FIELDING_E	0	0	0.0000000	0.000000	0	0.000000
FIELDING_DP	-286	0	1.0780691	3.000000	0	1.327290

Of all the observations gathered across these fifteen variables, 10.187% were missing information; batters hit by pitches was missing the most, with 2,085 instances of missing information. The standard deviation of the various variables hints at the intense skewing of some of the variables provided, especially the hits allowed, walks allowed and strike outs.

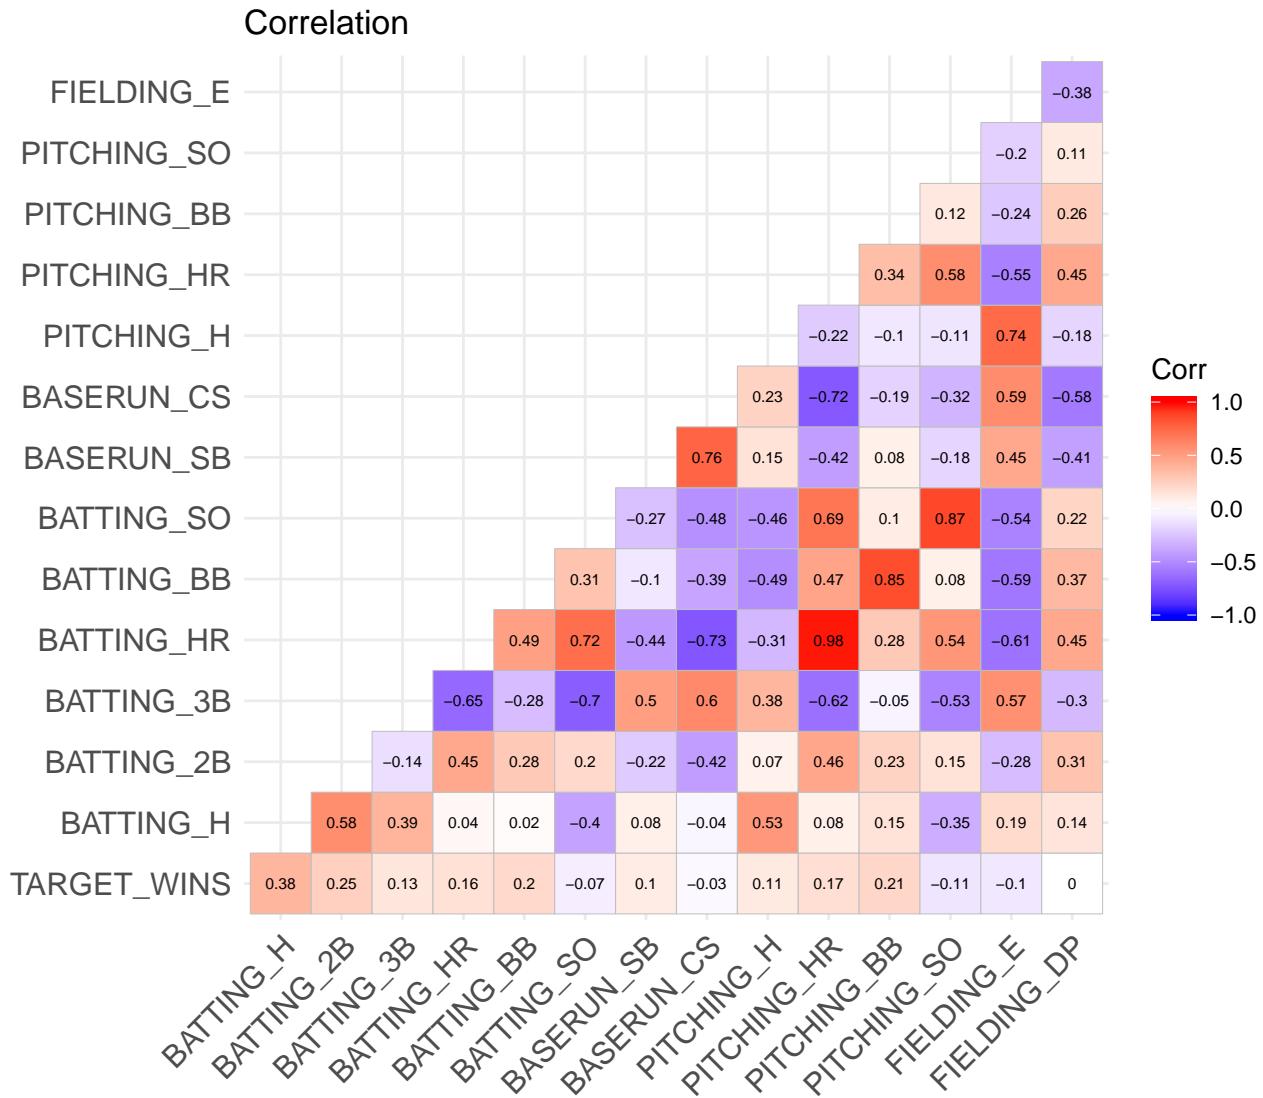
Histograms



Histograms



The theoretical effect of strikeouts by batters, batters caught stealing, errors, walks, hits, and homeruns allowed were believed to theoretically have a negative impact on the number of wins of an individual team in a given year. A closer look at the correlation between the variables painted a different picture.



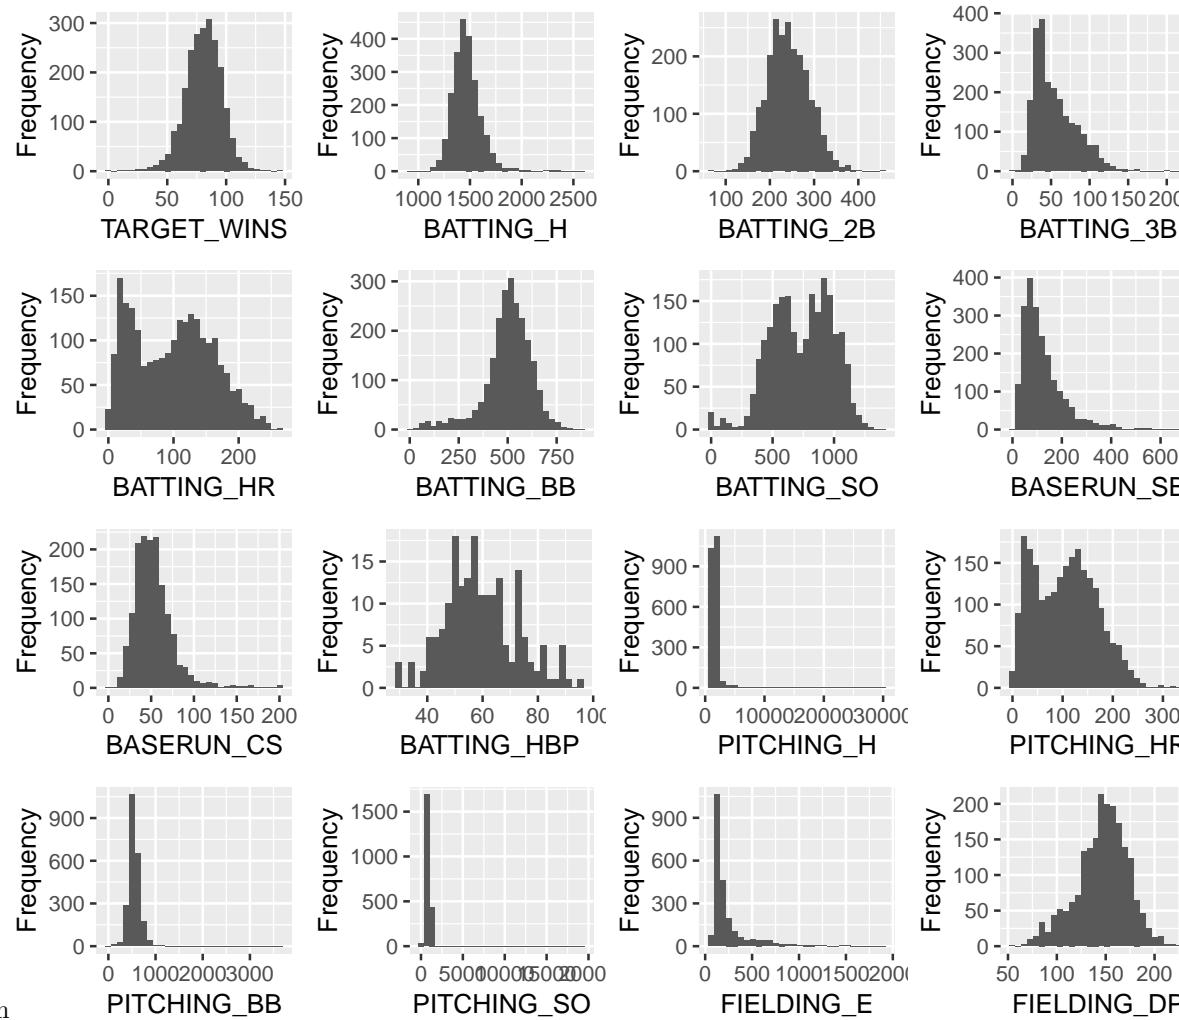
When compared to what was hypothesized, there was actually a positive impact for the number of wins for a team in a given year by walks, hits, and homeruns allowed; at the same time, variables previously thought to have a positive correlation - strikeouts by pitchers and double plays - had a negative correlation for the number of wins. The three variables with the greatest correlation to the number of wins were the hits allowed, the walks by batters, and the walks allowed. Of these, the hits allowed had a relatively low correlation with the walks by batters and the walks allowed, whereas the walks allowed and the walks by batters had a direct positive correlation with one another.

- Describe the size:

The money ball data is 144kb in size. The data contains 2,276 rows and 16 columns without the index. The variables are continuous integer. The TARGET_WINS is our response variable. There are 3,478 missing values out of 36,416 observations.

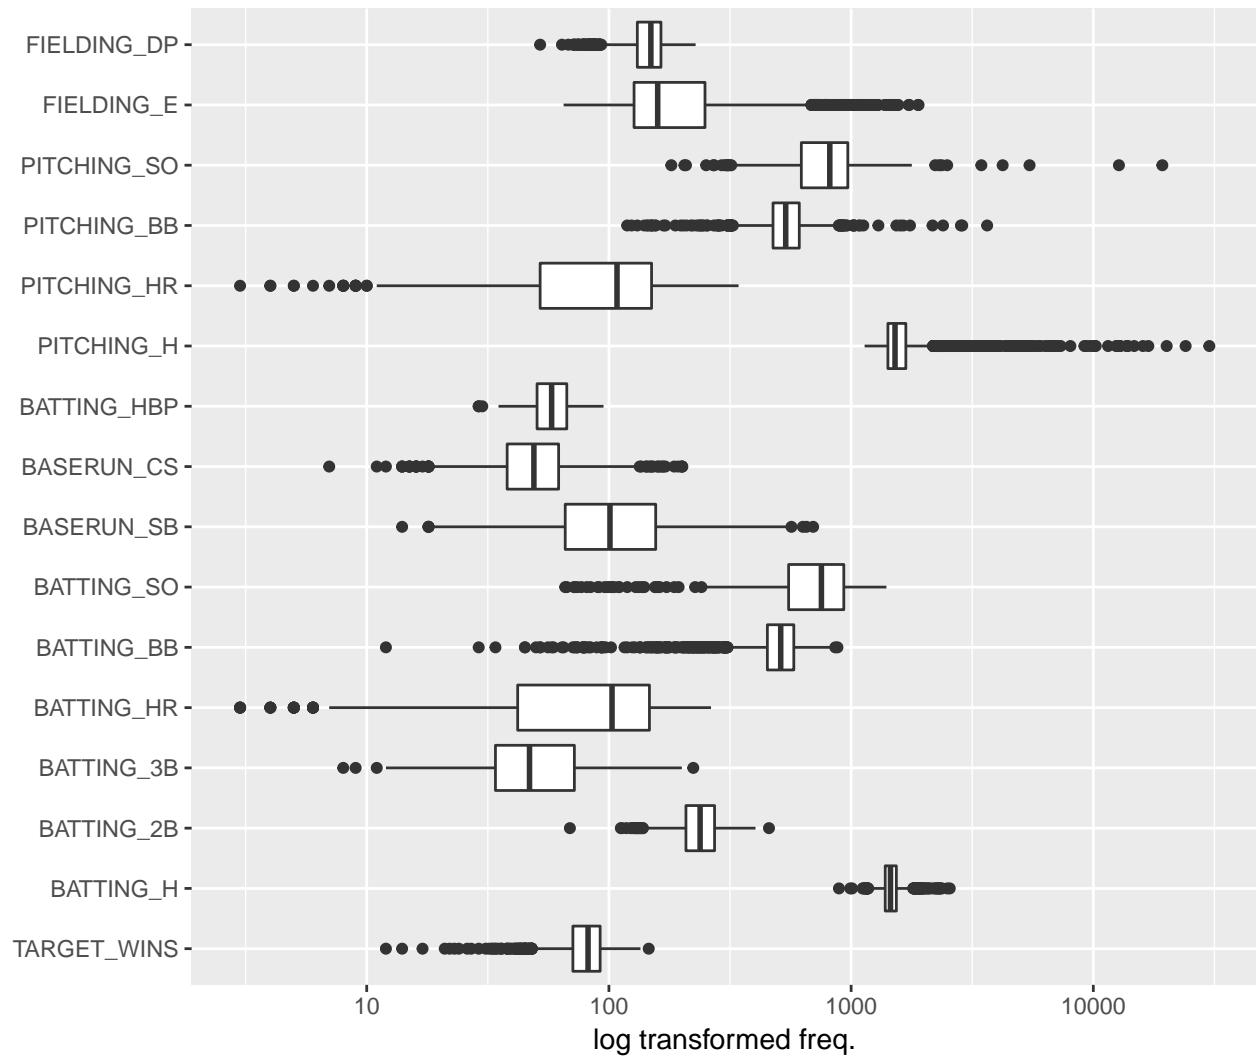
- Statistics summary

	vars	n	mean	sd	median	trimmed	mad	min	max	range
TARGET_WINS	1	2276	80.79086	15.75215	82.0	81.31229	14.8260	0	146	146
BATTING_H	2	2276	1469.26977	144.59120	1454.0	1459.04116	114.1602	891	2554	1663
BATTING_2B	3	2276	241.24692	46.80141	238.0	240.39627	47.4432	69	458	389
BATTING_3B	4	2276	55.25000	27.93856	47.0	52.17563	23.7216	0	223	223
BATTING_HR	5	2276	99.61204	60.54687	102.0	97.38529	78.5778	0	264	264
BATTING_BB	6	2276	501.55888	122.67086	512.0	512.18331	94.8864	0	878	878
BATTING_SO	7	2174	735.60534	248.52642	750.0	742.31322	284.6592	0	1399	1399
BASERUN_SB	8	2145	124.76177	87.79117	101.0	110.81188	60.7866	0	697	697
BASERUN_CS	9	1504	52.80386	22.95634	49.0	50.35963	17.7912	0	201	201
BATTING_HBP	10	191	59.35602	12.96712	58.0	58.86275	11.8608	29	95	66
PITCHING_H	11	2276	1779.21046	1406.84293	1518.0	1555.89517	174.9468	1137	30132	28995
PITCHING_HR	12	2276	105.69859	61.29875	107.0	103.15697	74.1300	0	343	343
PITCHING_BB	13	2276	553.00791	166.35736	536.5	542.62459	98.5929	0	3645	3645
PITCHING_SO	14	2174	817.73045	553.08503	813.5	796.93391	257.2311	0	19278	19278
FIELDING_E	15	2276	246.48067	227.77097	159.0	193.43798	62.2692	65	1898	1833
FIELDING_DP	16	1990	146.38794	26.22639	149.0	147.57789	23.7216	52	228	176

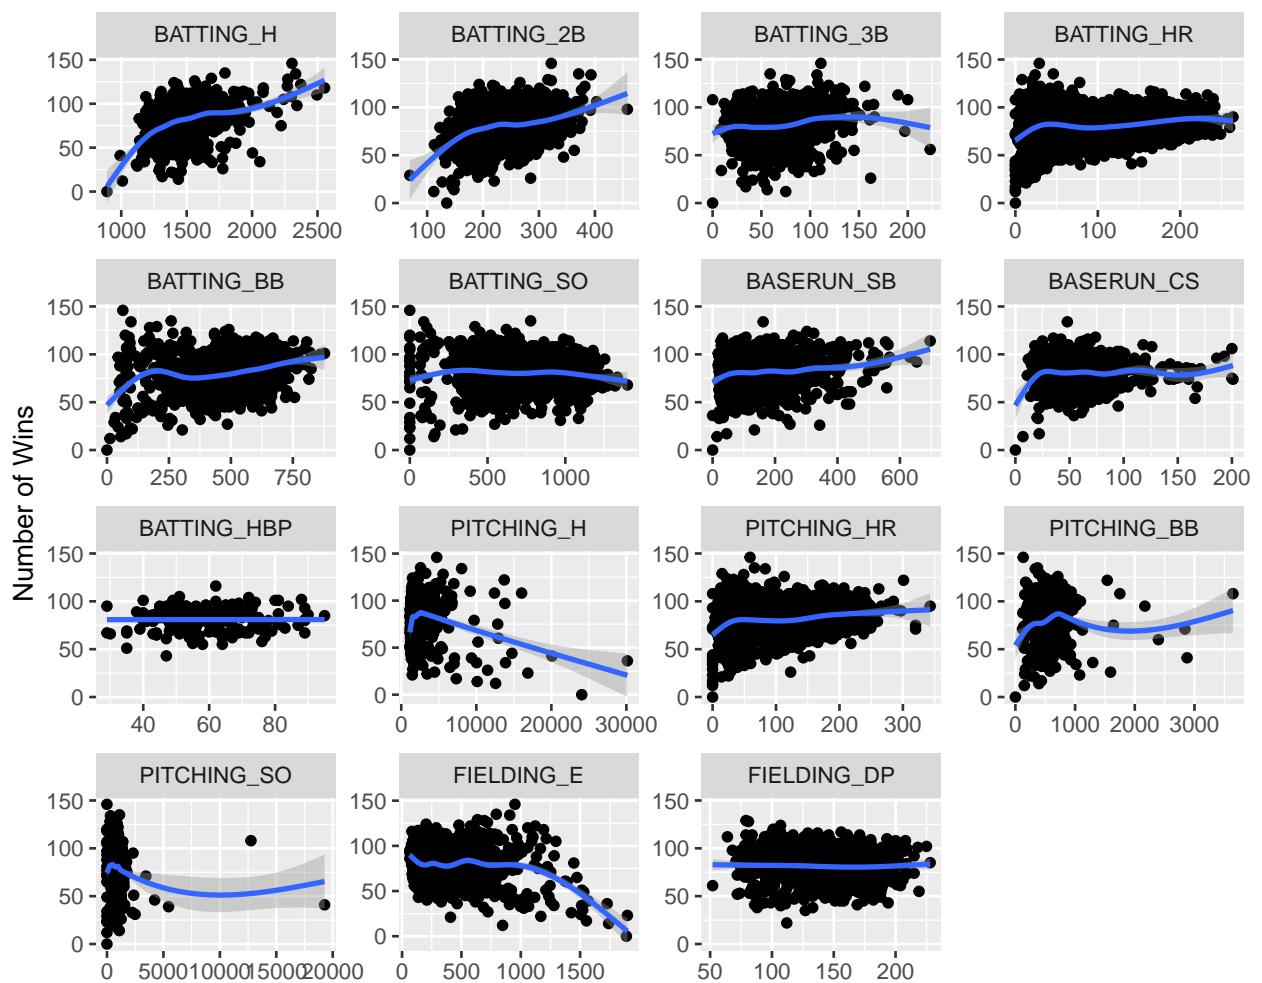


- Data visualization

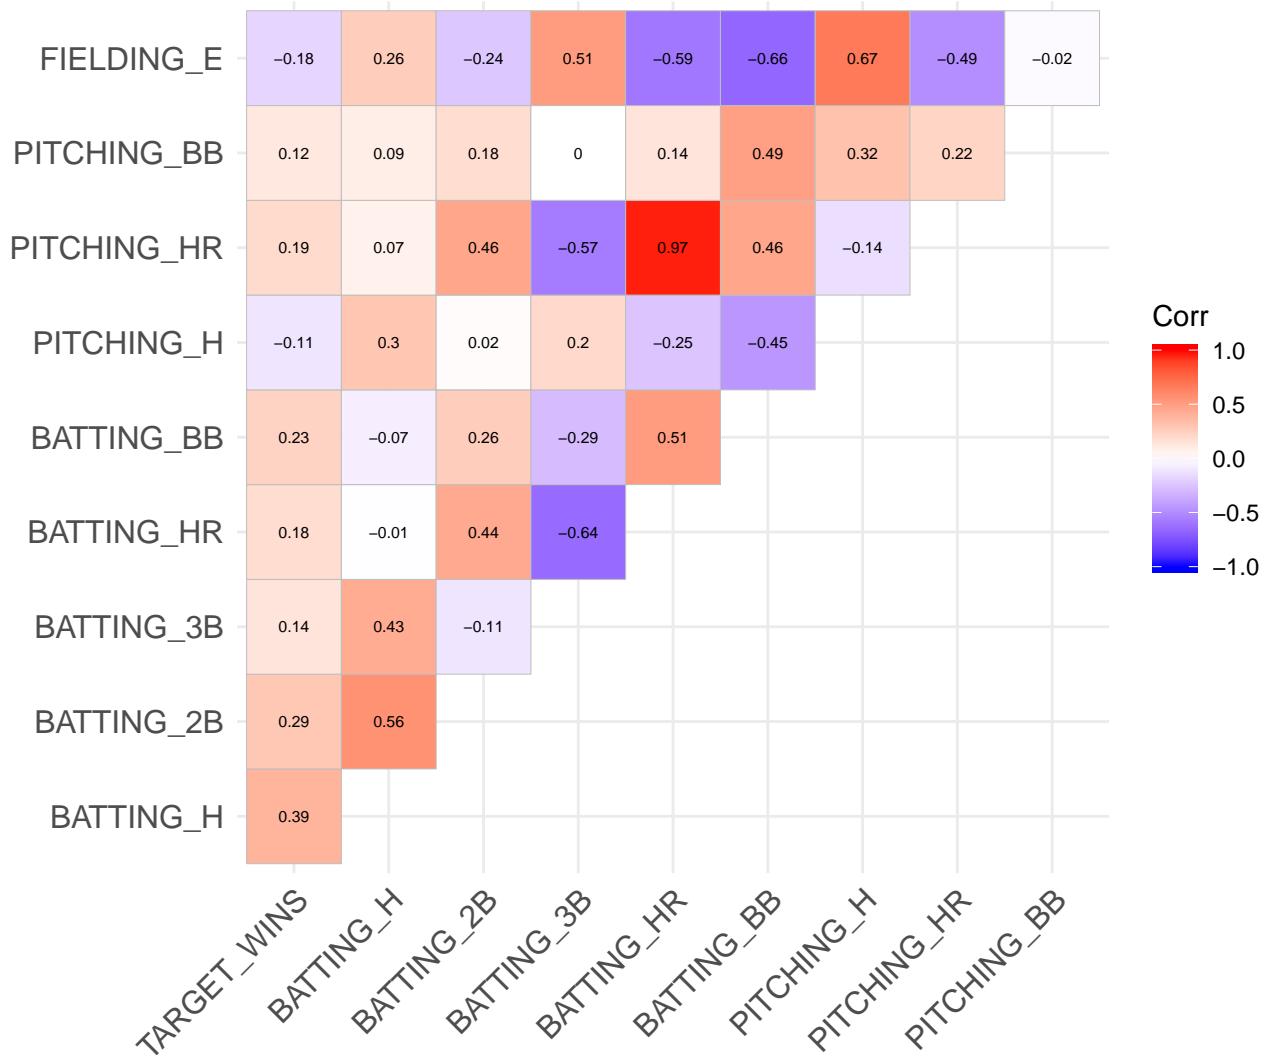
Boxplot

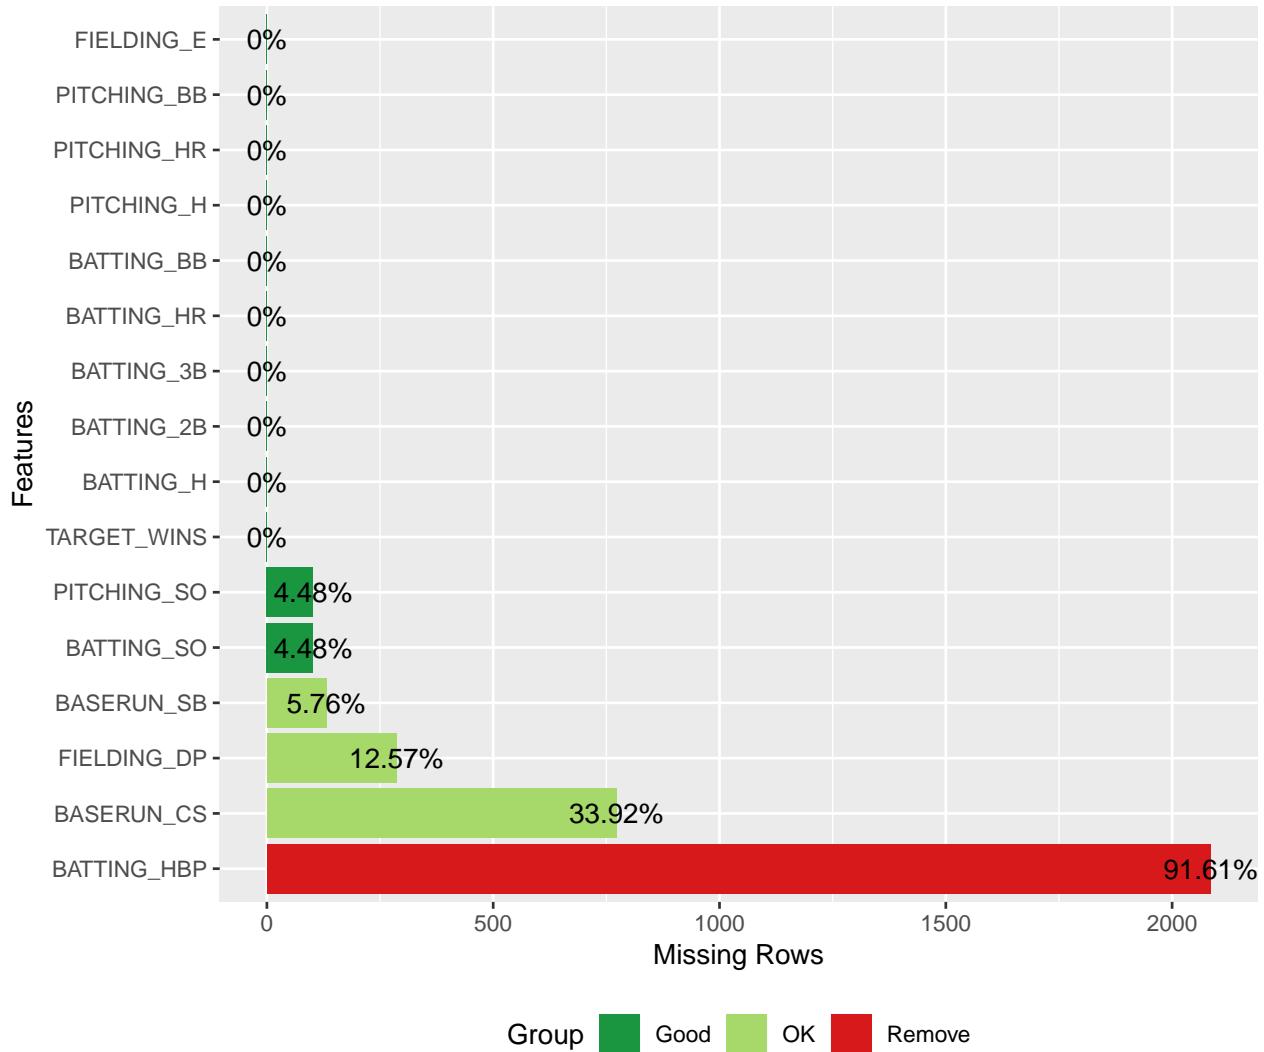


Scatterplot



Correlation





2 DATA PREPARATION

2.1 Instructions:

Describe how you have transformed the data by changing the original variables or creating new variables. If you did transform the data or create new variables, discuss why you did this. Here are some possible transformations.

- Fix missing values (maybe with a Mean or Median value)
- Create flags to suggest if a variable was missing
- Transform data by putting it into buckets
- Mathematical transforms such as log or square root
- Combine variables (such as ratios or adding or multiplying) to create new variables

2.2 Possible writeup for Data Preparation

As previously mentioned, just north of 10% of the data was missing values. Missing values can lead to errors in a model, bias, and worse if left unaccounted for. Attempting to “fix” this by imputing values or guessing

why the values are missing in the first place - such as concluding that the missing values are meant to be zeroes - are just as likely to help with creating a model as it is to help with creating a disaster.

One of the R packages utilized, DataExplorer, recommends removing null or missing values; it was for this reason all observations of hits by pitch were removed.<< SOURCE FOR THIS IS REQUIRED! IF NO SOURCE IS PROVIDED, CONSIDER USING “Due to the sheer volume of missing values present in the observations for hits by pitch (91.61%) it was determined the best course of action was to remove the variable altogether.” OR A VARIATION THEREOF. >> Deleting all cases with missing values, in this instance, would have shrunk the size of the dataset down to less than a tenth of its original size. For this reason, the feature itself was excluded from the dataset, rather than the cases that had no values present for it.

The other missing values - present in batting strikeouts... needs more work. x_x

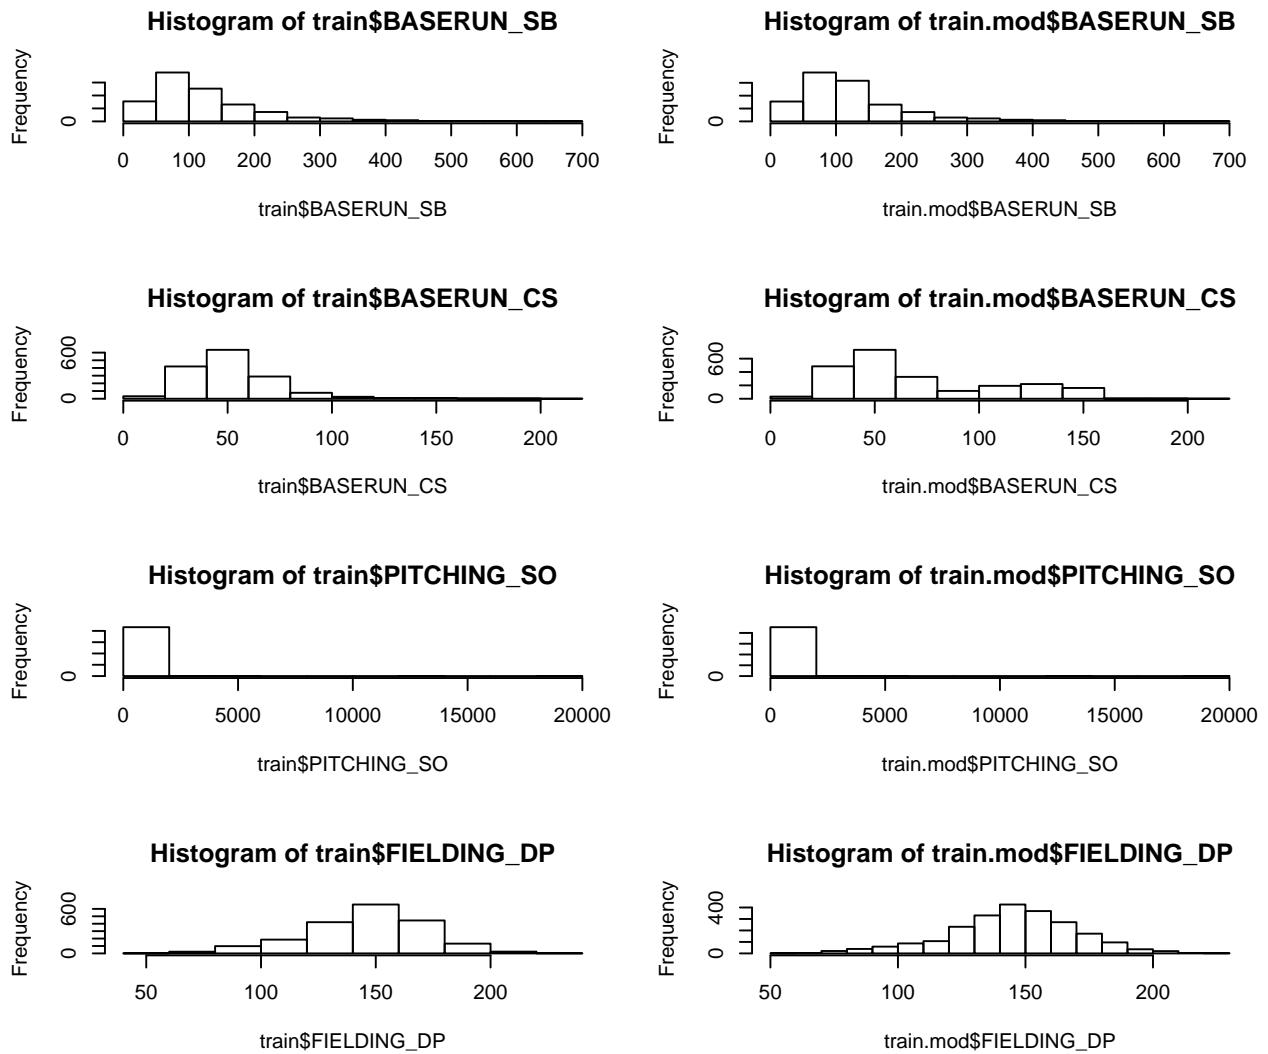
2.3 Information in general, don't use the writeups here as-is

2.4 Missing Values

- 1) Hit by pitch missing 91.61% .
 - Missing values can lead to errors and bias into a model. Fixing and imputation may help or make it worse.
 - When it is just a few observations missing, modifications can be made, however, with 91.61% is a large proportion and could distort the modelling later on that it is better to ignore this column.
 - The Data explorer package recommends to remove.
 - From LMR: Missing Completely at Random (MCAR) The probability that a value is missing is the same for all cases. If we simply delete all cases with missing values from the analysis, we will cause no bias, although we may lose some information.
 - However, there is no consensus on when to exclude missing data. Some argue that missing data more than 10% can lead to bias. Others argue that missing data patterns have greater impact than the proportion.
- 2) Pitching_S0 and Batting_S0 are missing exact same proportion 4.48% and are missing in the same observations.

2.5 NA Imputation

```
## TARGET_WINS    BATTING_H    BATTING_2B    BATTING_3B    BATTING_HR    BATTING_BB
##          0          0          0          0          0          0
##  BATTING_S0    BASERUN_SB   BASERUN_CS    PITCHING_H    PITCHING_HR   PITCHING_BB
##         102        131        772          0          0          0
##  PITCHING_S0    FIELDING_E   FIELDING_DP
##         102          0        286
```



2.6 Feature Engineering

Jeremy: Adjusted this to reflect offense (batting) minus defense (pitching). These arithmetically transformed offense / defense variables are linearly related with BATTING and PITCHING variables, so we can include one or the other in a model, but not both. Replacing original variables with these transforms did not improve R^2 in a base case.

FOR THE OTHER HALF OF THE GROUP:

```
z_train <- sapply(imputed_train, scale)
log_train <- log(imputed_train) # weird results
z_log_train <- sapply(log_train, scale) # weirder results
```

`imputed_train` is most likely the variable you want to use.

3 BUILD MODELS

3.1 Instructions:

Using the training data set, build at least three different multiple linear regression models, using different variables (or the same variables with different transformations). Since we have not yet covered automated variable selection methods, you should select the variables manually (unless you previously learned Forward or Stepwise selection, etc.). Since you manually selected a variable for inclusion into the model or exclusion into the model, indicate why this was done. Discuss the coefficients in the models, do they make sense? For example, if a team hits a lot of Home Runs, it would be reasonably expected that such a team would win more games. However, if the coefficient is negative (suggesting that the team would lose more games), then that needs to be discussed. Are you keeping the model even though it is counter intuitive? Why? The boss needs to know.

3.2 MODEL 1

Multiple regression can be created as purely statistical model, through the use of significance tests, or it can be interpreted in a more practical, non-statistical manner. This approach based on the subject-area expertise.

We've created following categories from the most important to the least important variables according to the subject-area expert.

Very Important: BATTING_H, BATTING_HR, BATTING_SO ,FIELDING_E, PITCHING_SO

Fairly Important: BASERUN_SB, PITCHING_HR, BATTING_BB

Important: BATTING_2B, BATTING_3B, FIELDING_DP, PITCHING_H

Slightly Important: PITCHING_BB, BASERUN_CS

Not at all important: BATTING_HBP

'Batters hit by pitch' and 'Caught Stealing' have been eliminated as least important variables.

```
##  
## Call:  
## lm(formula = TARGET_WINS ~ BATTING_H + BATTING_HR + BATTING_SO +  
##      FIELDING_E + PITCHING_SO + BASERUN_SB + PITCHING_HR + BATTING_BB +  
##      BATTING_2B + BATTING_3B + FIELDING_DP + PITCHING_BB + PITCHING_H,  
##      data = imputed_train)  
##  
## Residuals:  
##      Min        1Q    Median        3Q       Max  
## -81.602   -8.272    0.065    7.985   69.047  
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 36.592272  5.663910  6.461 1.28e-10 ***  
## BATTING_H    0.016719  0.005010  3.337 0.000861 ***  
## BATTING_HR   0.119188  0.050659  2.353 0.018723 *  
## BATTING_SO   -0.031132  0.006598 -4.718 2.53e-06 ***  
## FIELDING_E   -0.038259  0.003433 -11.144 < 2e-16 ***  
## PITCHING_SO  0.019966  0.005420  3.684 0.000235 ***  
## BASERUN_SB   0.041446  0.004934  8.399 < 2e-16 ***  
## PITCHING_HR  -0.034985  0.046876 -0.746 0.455549  
## BATTING_BB    0.078335  0.016217  4.830 1.46e-06 ***
```

```

## BATTING_2B -0.011063 0.009390 -1.178 0.238836
## BATTING_3B 0.126375 0.018155 6.961 4.43e-12 ***
## FIELDING_DP -0.091226 0.013111 -6.958 4.53e-12 ***
## PITCHING_BB -0.055706 0.014380 -3.874 0.000110 ***
## PITCHING_H 0.013513 0.001539 8.779 < 2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.63 on 2216 degrees of freedom
## Multiple R-squared: 0.2835, Adjusted R-squared: 0.2793
## F-statistic: 67.46 on 13 and 2216 DF, p-value: < 2.2e-16

## Start: AIC=11325.02
## TARGET_WINS ~ BATTING_H + BATTING_HR + BATTING_SO + FIELDING_E +
##      PITCHING_SO + BASERUN_SB + PITCHING_HR + BATTING_BB + BATTING_2B +
##      BATTING_3B + FIELDING_DP + PITCHING_BB + PITCHING_H
##
##          Df Sum of Sq   RSS   AIC
## - PITCHING_HR 1     88.9 353605 11324
## - BATTING_2B   1    221.5 353737 11324
## <none>                   353516 11325
## - BATTING_HR   1    883.1 354399 11329
## - BATTING_H    1   1776.5 355292 11334
## - PITCHING_SO  1   2165.1 355681 11337
## - PITCHING_BB  1   2393.9 355910 11338
## - BATTING_SO   1   3551.6 357068 11345
## - BATTING_BB   1   3722.3 357238 11346
## - FIELDING_DP  1   7723.5 361239 11371
## - BATTING_3B   1   7730.1 361246 11371
## - BASERUN_SB   1  11255.0 364771 11393
## - PITCHING_H   1  12294.0 365810 11399
## - FIELDING_E   1  19811.3 373327 11445
##
## Step: AIC=11323.58
## TARGET_WINS ~ BATTING_H + BATTING_HR + BATTING_SO + FIELDING_E +
##      PITCHING_SO + BASERUN_SB + BATTING_BB + BATTING_2B + BATTING_3B +
##      FIELDING_DP + PITCHING_BB + PITCHING_H
##
##          Df Sum of Sq   RSS   AIC
## - BATTING_2B   1   218.7 353824 11323
## <none>                   353605 11324
## - BATTING_H    1   1790.1 355395 11333
## - PITCHING_SO  1   2131.4 355736 11335
## - BATTING_SO   1   3491.5 357096 11344
## - PITCHING_BB  1   4558.7 358164 11350
## - BATTING_BB   1   6449.9 360055 11362
## - BATTING_3B   1   7663.0 361268 11369
## - FIELDING_DP  1   7822.8 361428 11370
## - BATTING_HR   1   11191.6 364796 11391
## - BASERUN_SB   1   11302.9 364908 11392
## - PITCHING_H   1   12384.3 365989 11398
## - FIELDING_E   1   20640.3 374245 11448
##
## Step: AIC=11322.96

```

```

## TARGET_WINS ~ BATTING_H + BATTING_HR + BATTING_SO + FIELDING_E +
##      PITCHING_SO + BASERUN_SB + BATTING_BB + BATTING_3B + FIELDING_DP +
##      PITCHING_BB + PITCHING_H
##
##          Df Sum of Sq    RSS   AIC
## <none>            353824 11323
## - BATTING_H     1    1777.1 355601 11332
## - PITCHING_SO  1    1990.8 355814 11334
## - BATTING_SO   1    3479.4 357303 11343
## - PITCHING_BB  1    4495.2 358319 11349
## - BATTING_BB   1    6351.2 360175 11361
## - FIELDING_DP  1    7773.8 361597 11369
## - BATTING_3B   1    8094.2 361918 11371
## - BATTING_HR   1   11452.8 365276 11392
## - BASERUN_SB   1   11828.7 365652 11394
## - PITCHING_H   1   12901.1 366725 11401
## - FIELDING_E   1   20437.9 374261 11446

##
## Call:
## lm(formula = TARGET_WINS ~ BATTING_H + BATTING_HR + BATTING_SO +
##      FIELDING_E + PITCHING_SO + BASERUN_SB + BATTING_BB + BATTING_3B +
##      FIELDING_DP + PITCHING_BB + PITCHING_H, data = imputed_train)
##
## Coefficients:
## (Intercept)    BATTING_H    BATTING_HR    BATTING_SO    FIELDING_E
## 38.93008       0.01316      0.08286     -0.03074     -0.03727
## PITCHING_SO    BASERUN_SB   BATTING_BB   BATTING_3B   FIELDING_DP
## 0.01896        0.04219      0.08441      0.12819     -0.09140
## PITCHING_BB    PITCHING_H
## -0.06157       0.01375

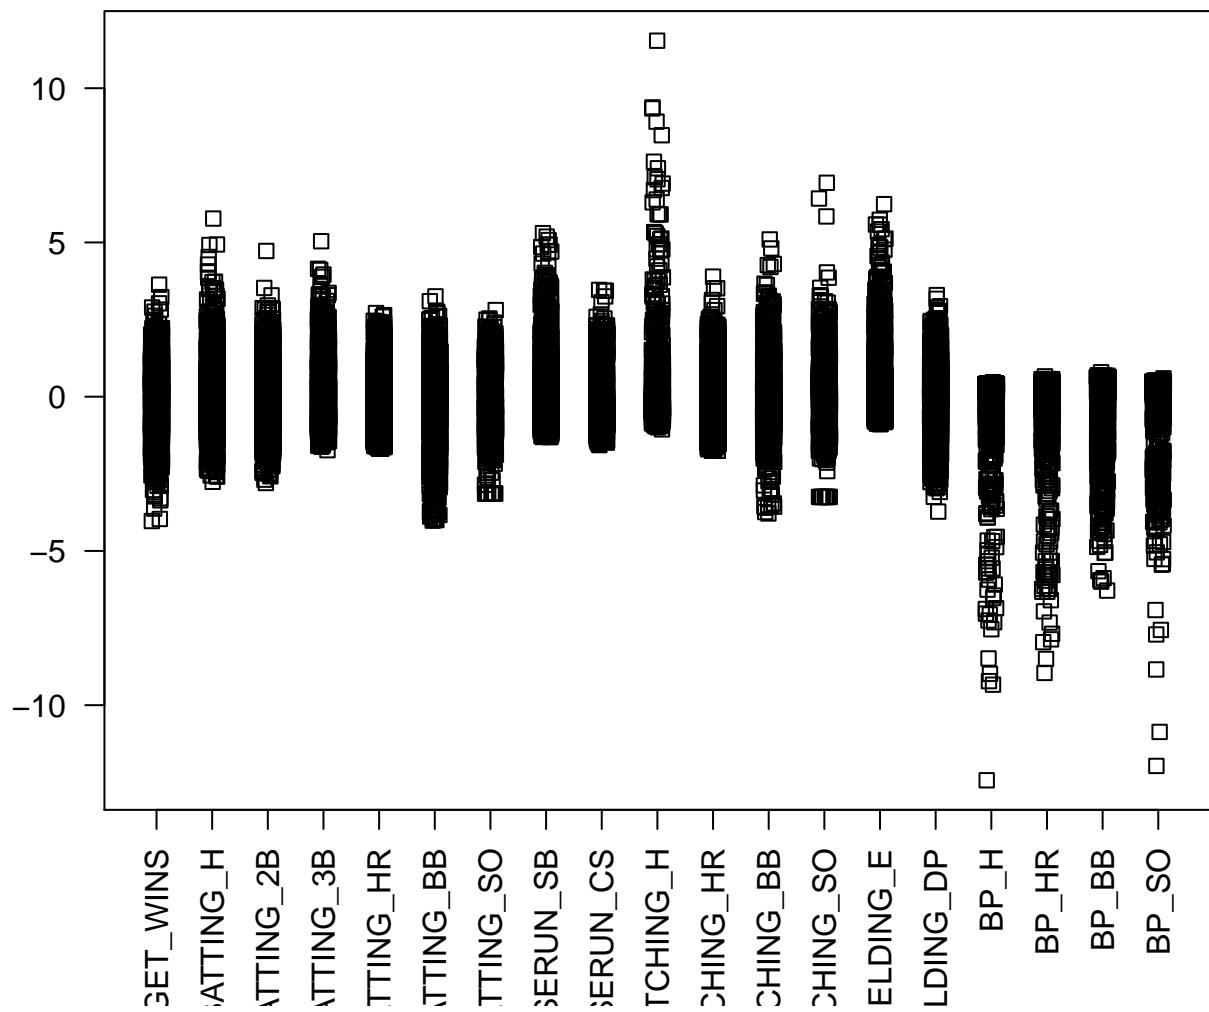
##
## Call:
## lm(formula = TARGET_WINS ~ BATTING_HR + BATTING_SO + FIELDING_E +
##      PITCHING_SO + BASERUN_SB + PITCHING_HR + BATTING_BB + BATTING_3B +
##      FIELDING_DP + PITCHING_BB + PITCHING_H, data = imputed_train)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -92.349 -8.319  0.095  8.044  73.821
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 50.967361  3.815685 13.357 < 2e-16 ***
## BATTING_HR  0.137603  0.050478  2.726 0.006461 **
## BATTING_SO -0.033187  0.006582 -5.042 4.98e-07 ***
## FIELDING_E -0.041649  0.003245 -12.835 < 2e-16 ***
## PITCHING_SO 0.018692  0.005384  3.472 0.000527 ***
## BASERUN_SB  0.045791  0.004785  9.569 < 2e-16 ***
## PITCHING_HR -0.038873  0.046964 -0.828 0.407915
## BATTING_BB   0.085893  0.016055  5.350 9.70e-08 ***
## BATTING_3B   0.156437  0.016039  9.754 < 2e-16 ***
## FIELDING_DP -0.083573  0.012950 -6.454 1.34e-10 ***
## PITCHING_BB -0.061960  0.014270 -4.342 1.48e-05 ***

```

```

## PITCHING_H 0.016931 0.001185 14.286 < 2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.66 on 2218 degrees of freedom
## Multiple R-squared: 0.2795, Adjusted R-squared: 0.276
## F-statistic: 78.23 on 11 and 2218 DF, p-value: < 2.2e-16

```



```

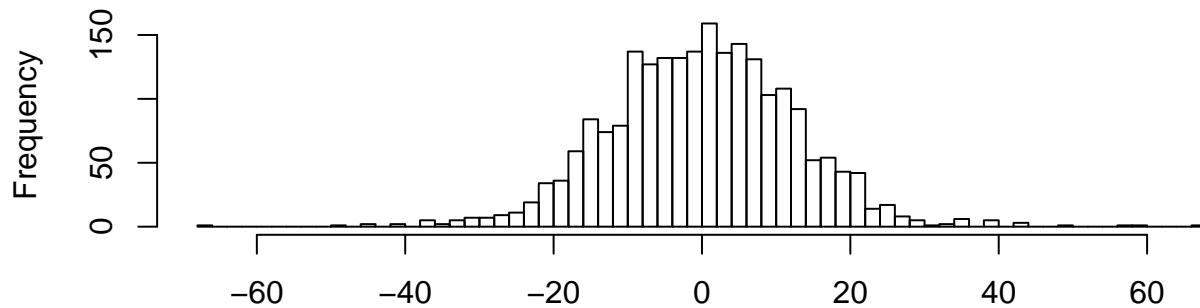
##
## Call:
## lm(formula = TARGET_WINS ~ BATTING_HR + BATTING_SO + FIELDING_E +
##     PITCHING_SO + BASERUN_SB + PITCHING_HR + BATTING_BB + BATTING_3B +
##     FIELDING_DP + PITCHING_BB + log(PITCHING_H), data = imputed_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -67.238  -8.139   0.185   8.020  67.777 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -3.037e+02  2.426e+01 -12.522 < 2e-16 ***
## BATTING_HR   1.134e-01  5.003e-02   2.268   0.0234 *  
## 
```

```

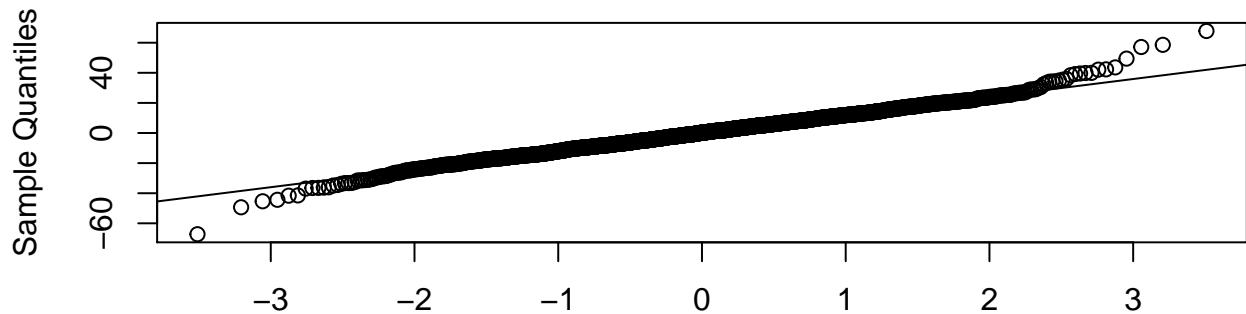
## BATTING_SO      -1.544e-02  6.668e-03  -2.315   0.0207 *
## FIELDING_E     -4.202e-02  3.159e-03  -13.299  < 2e-16 ***
## PITCHING_SO    7.688e-03  5.336e-03   1.441   0.1498
## BASERUN_SB     4.138e-02  4.680e-03   8.842  < 2e-16 ***
## PITCHING_HR    -4.352e-02  4.640e-02  -0.938   0.3484
## BATTING_BB      1.011e-01  1.600e-02   6.319  3.16e-10 ***
## BATTING_3B      1.171e-01  1.594e-02   7.349  2.79e-13 ***
## FIELDING_DP    -9.430e-02  1.283e-02  -7.351  2.76e-13 ***
## PITCHING_BB    -7.559e-02  1.423e-02  -5.311  1.20e-07 ***
## log(PITCHING_H) 5.222e+01  3.243e+00  16.099  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.52 on 2218 degrees of freedom
## Multiple R-squared:  0.2956, Adjusted R-squared:  0.2921
## F-statistic:  84.6 on 11 and 2218 DF,  p-value: < 2.2e-16

```

Histogram of Residuals



Normal Q-Q Plot



Theoretical Quantiles

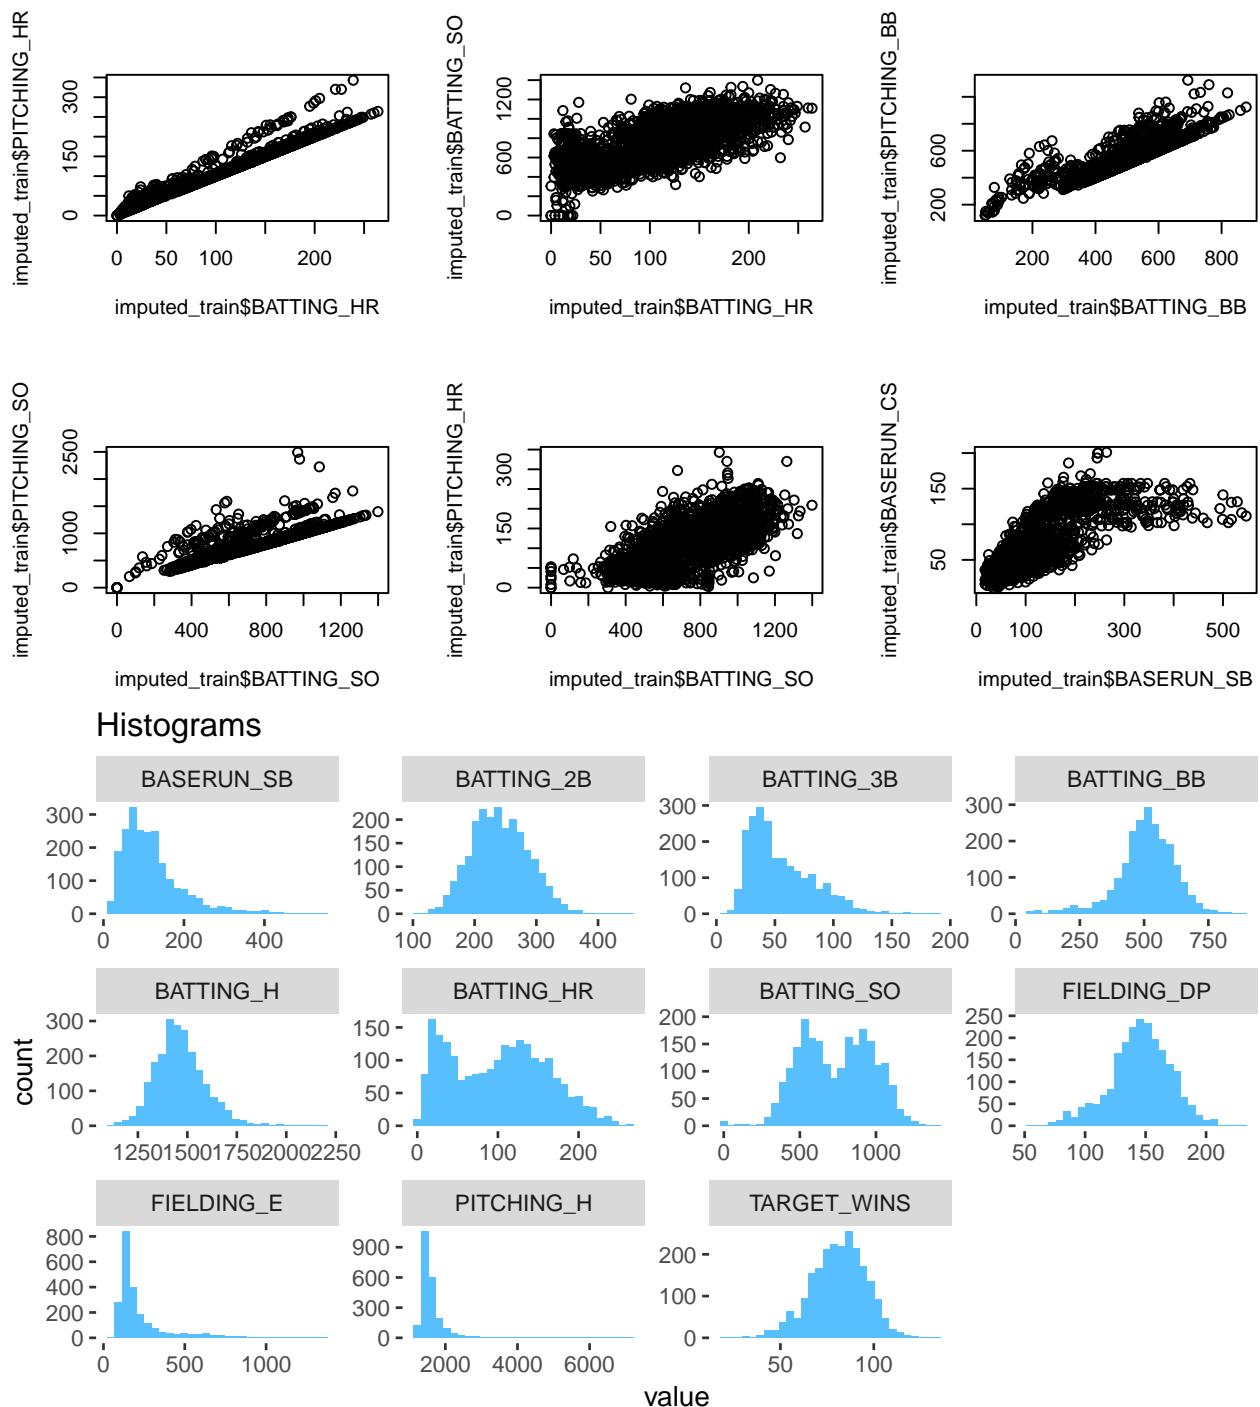
	r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df
value	0.2955663	0.2920727	12.51838	84.6026	0	12	-8793.869	17613.74	17687.96	347582.5	

```

##      Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## 60.82 76.09 81.12 81.44 86.10 108.26 54

```

3.3 MODEL 2



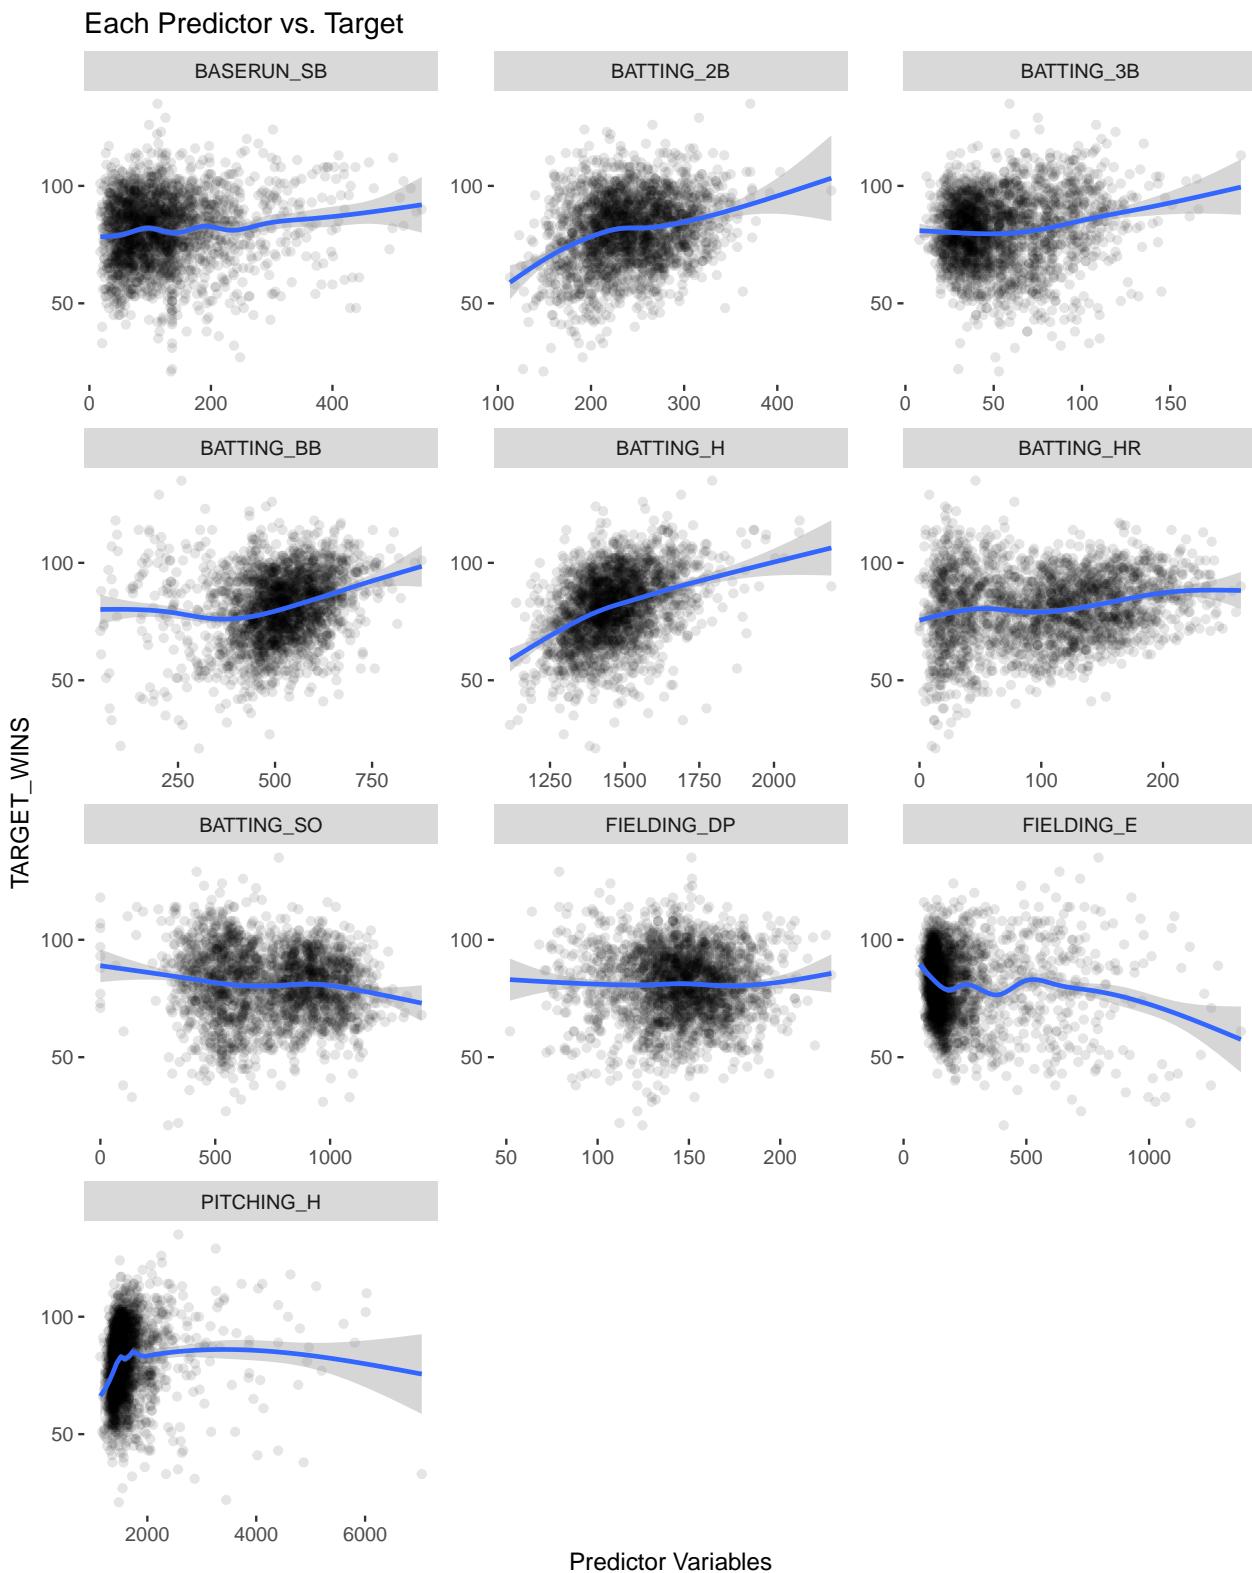
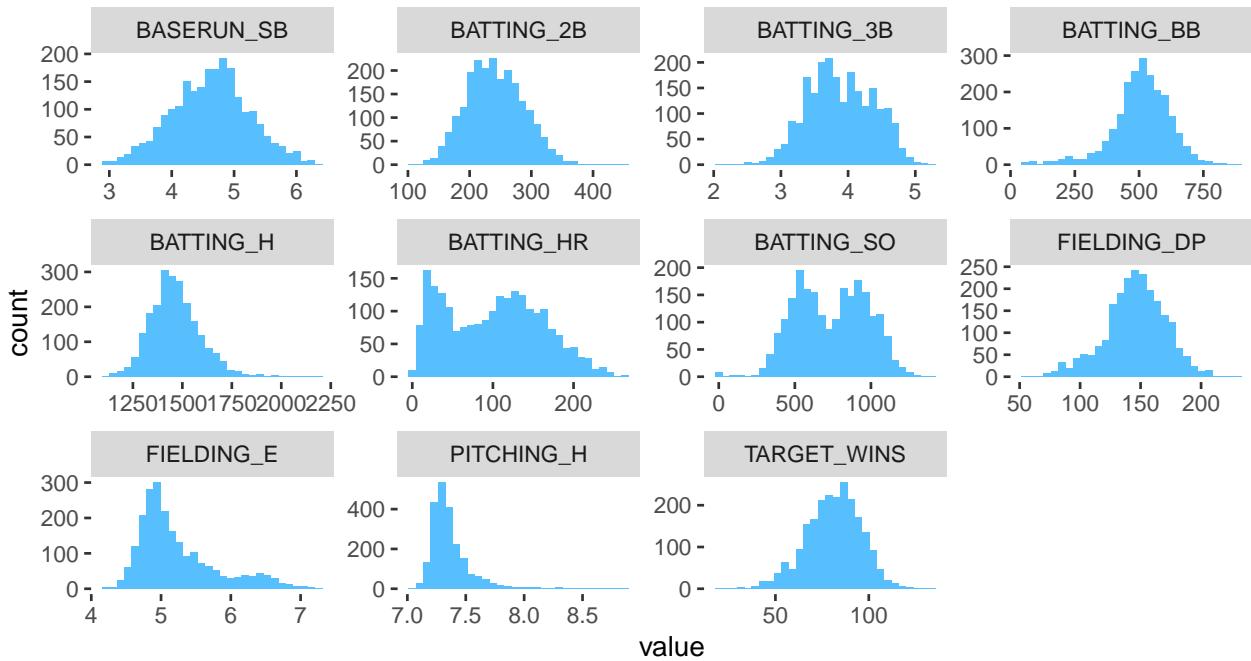


Figure 1: Each Predictor vs. Target

Histograms



```
##
## Call:
## lm(formula = TARGET_WINS ~ ., data = lm_data)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -52.645 -8.088  0.030   8.174  67.177 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -62.972452  15.913904 -3.957 7.83e-05 ***
## BATTING_H     0.025562  0.004611  5.543 3.32e-08 ***
## BATTING_2B   -0.029760  0.009216 -3.229  0.00126 ** 
## BATTING_3B    6.969781  0.947655  7.355 2.68e-13 ***
## BATTING_HR    0.070328  0.010126  6.946 4.93e-12 ***
## BATTING_BB    0.019229  0.003187  6.033 1.88e-09 ***
## BATTING_SO   -0.011228  0.002411 -4.656 3.41e-06 ***
## BASERUN_SB    4.546008  0.561766  8.092 9.54e-16 ***
## PITCHING_H   19.140858  2.616154  7.316 3.54e-13 ***
## FIELDING_E   -13.102762  1.061861 -12.339 < 2e-16 ***
## FIELDING_DP   -0.106723  0.013138 -8.123 7.47e-16 ***
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 12.57 on 2219 degrees of freedom
## Multiple R-squared:  0.2889, Adjusted R-squared:  0.2857 
## F-statistic: 90.14 on 10 and 2219 DF,  p-value: < 2.2e-16
```

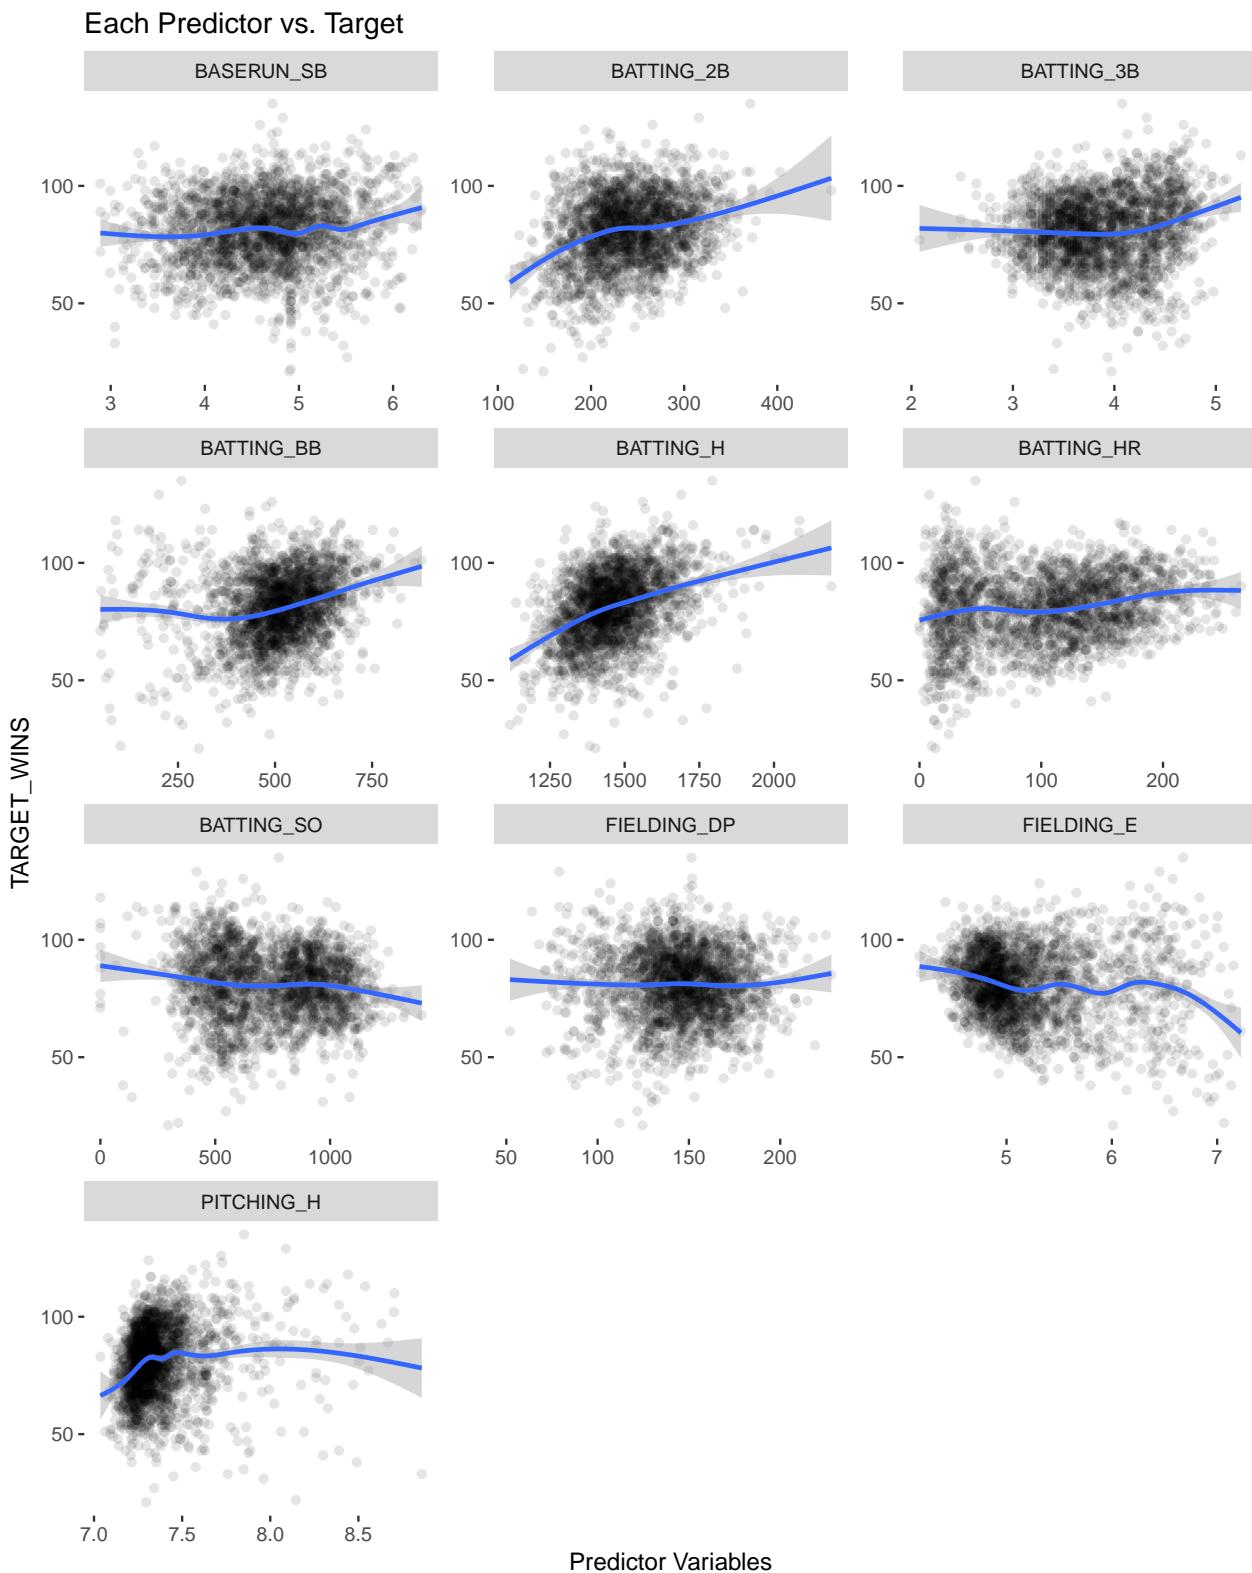
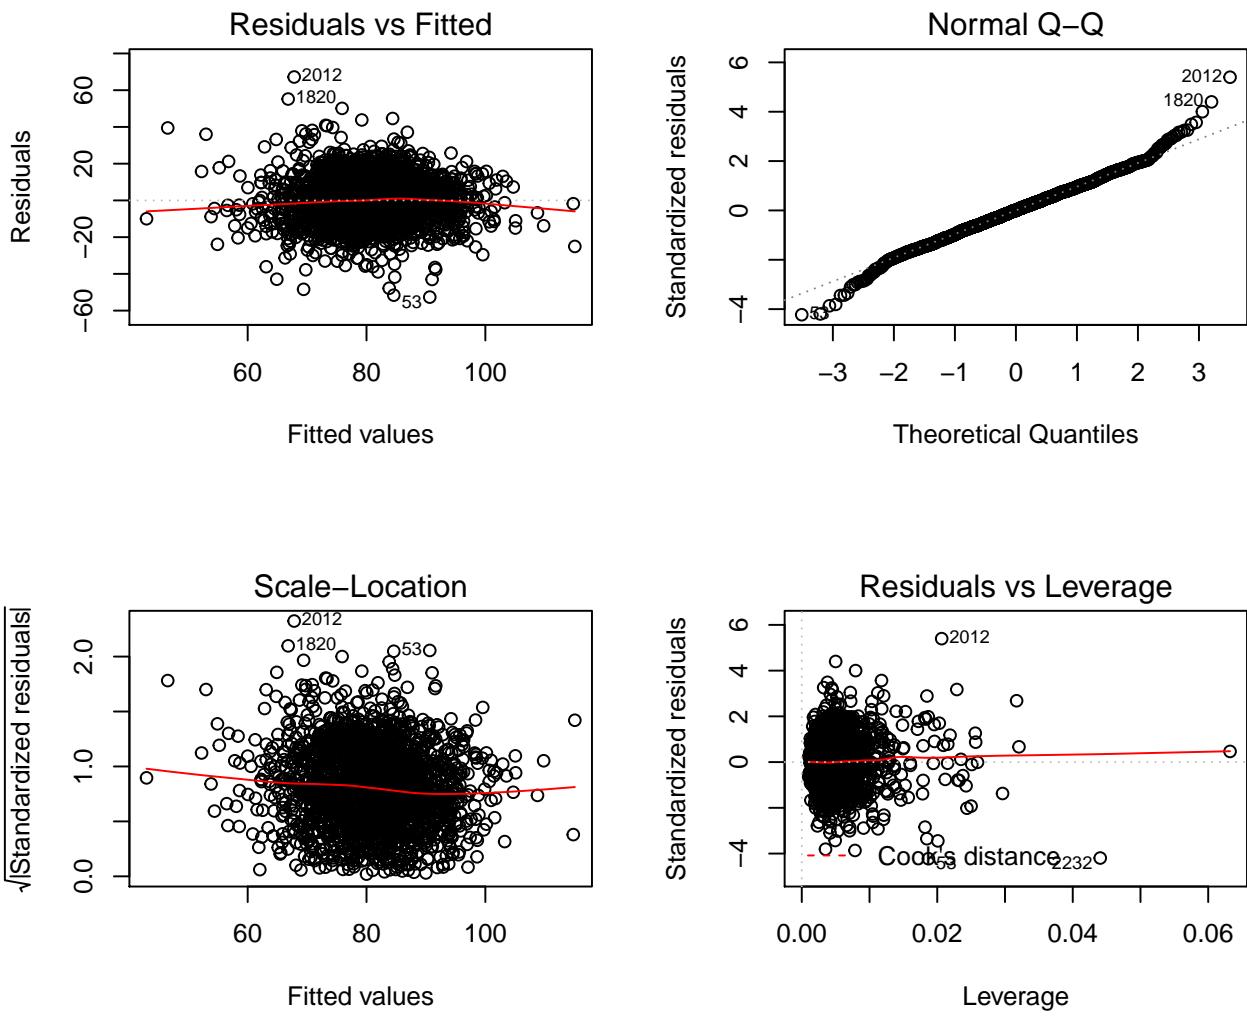


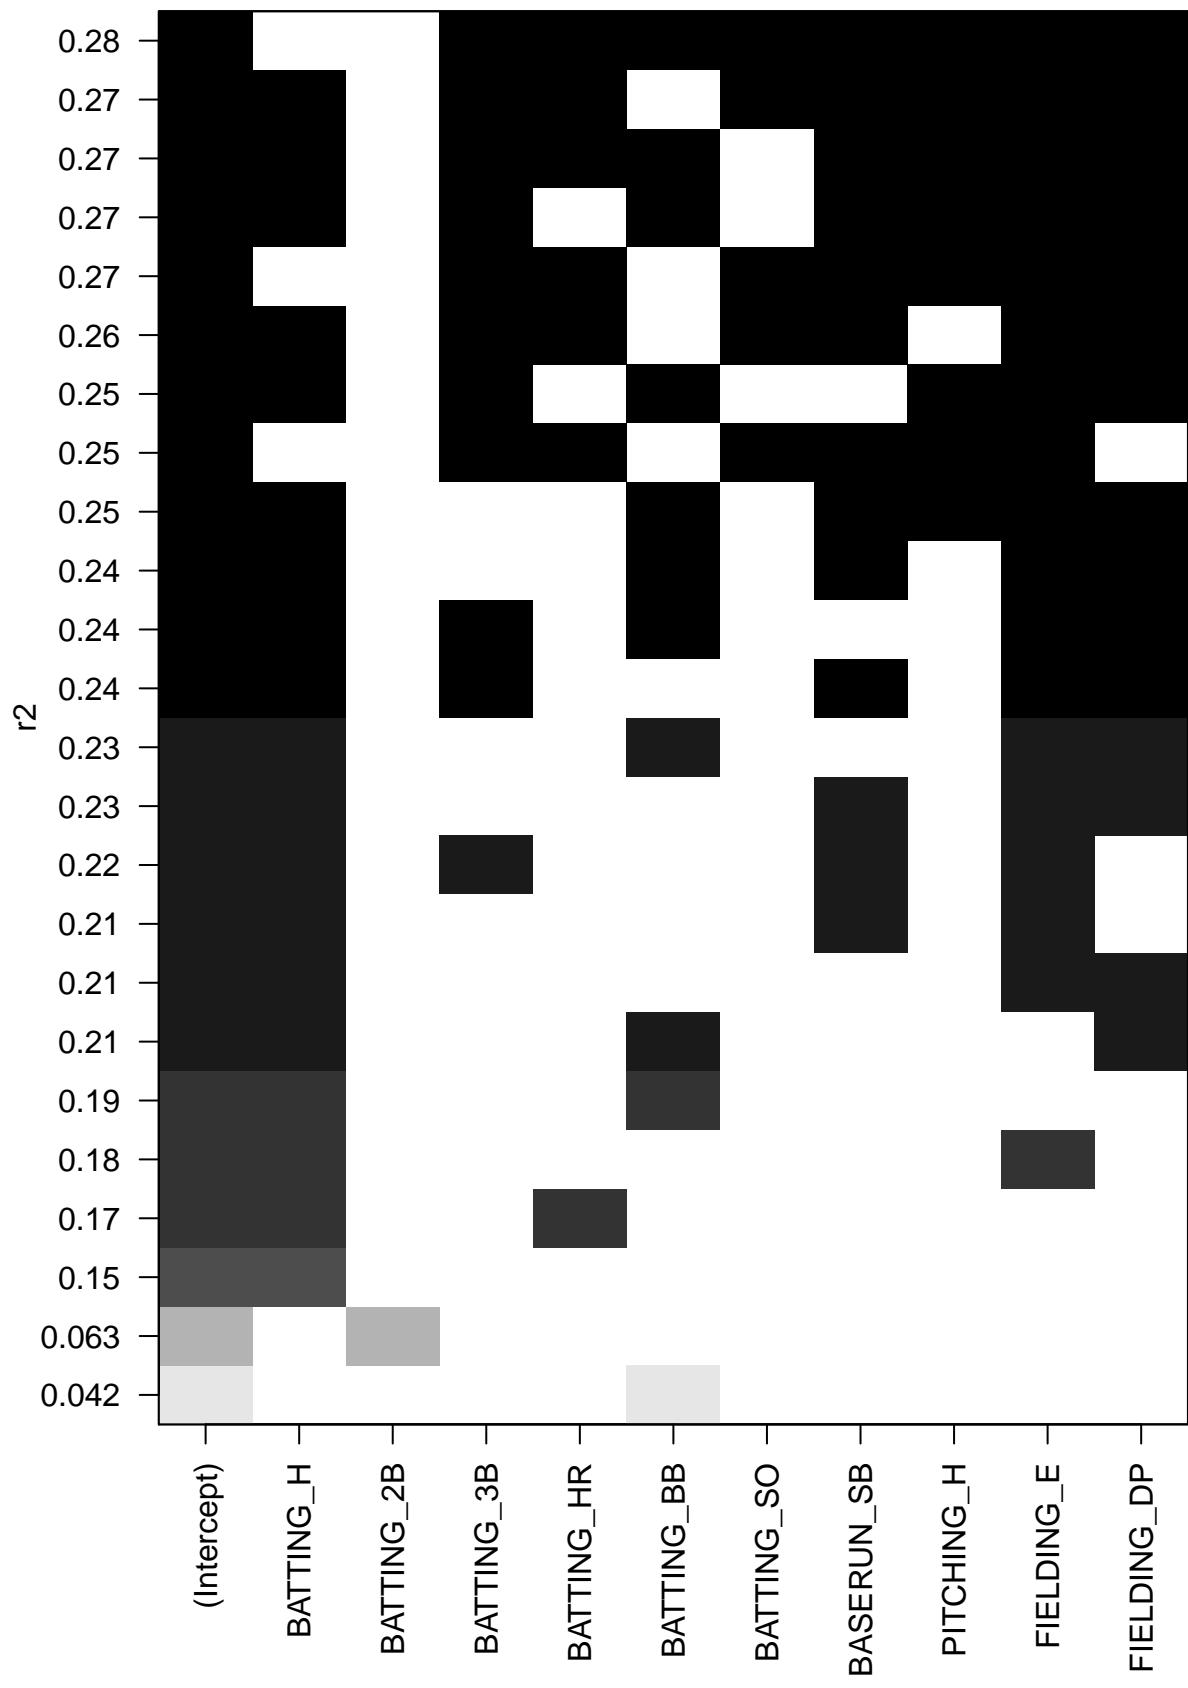
Figure 2: Each Predictor vs. Target

Table 1: Full Model Coefficients

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-62.9724524	15.9139042	-3.957071	0.0000783
BATTING_H	0.0255615	0.0046111	5.543426	0.0000000
BATTING_2B	-0.0297598	0.0092155	-3.229309	0.0012590
BATTING_3B	6.9697812	0.9476551	7.354765	0.0000000
BATTING_HR	0.0703282	0.0101256	6.945579	0.0000000
BATTING_BB	0.0192290	0.0031872	6.033192	0.0000000
BATTING_SO	-0.0112281	0.0024113	-4.656450	0.0000034
BASERUN_SB	4.5460080	0.5617657	8.092356	0.0000000
PITCHING_H	19.1408576	2.6161541	7.316411	0.0000000
FIELDING_E	-13.1027616	1.0618606	-12.339437	0.0000000
FIELDING_DP	-0.1067225	0.0131384	-8.122924	0.0000000



```
## NULL
```

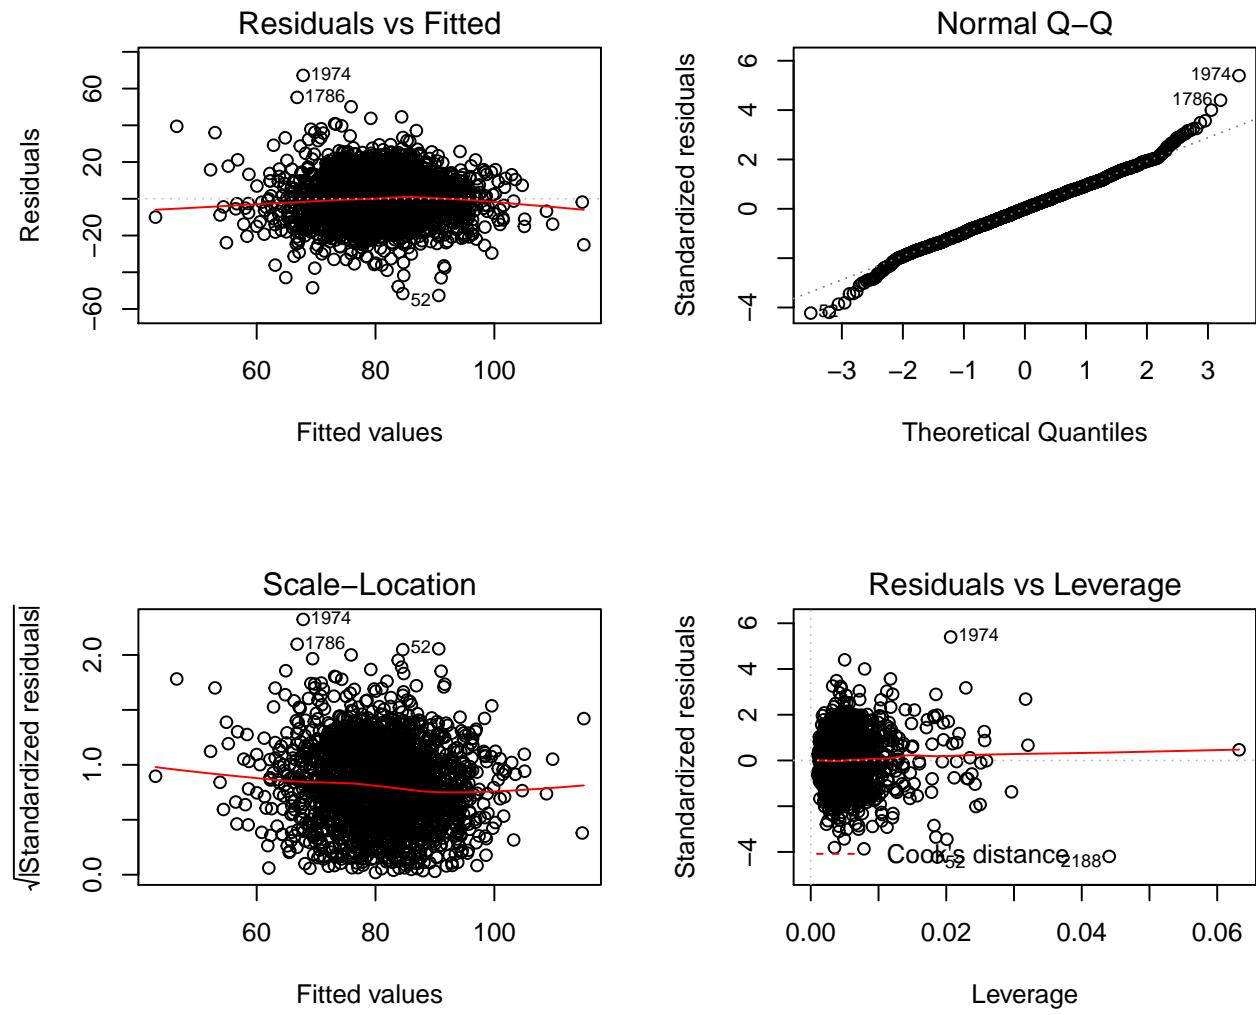


NULL

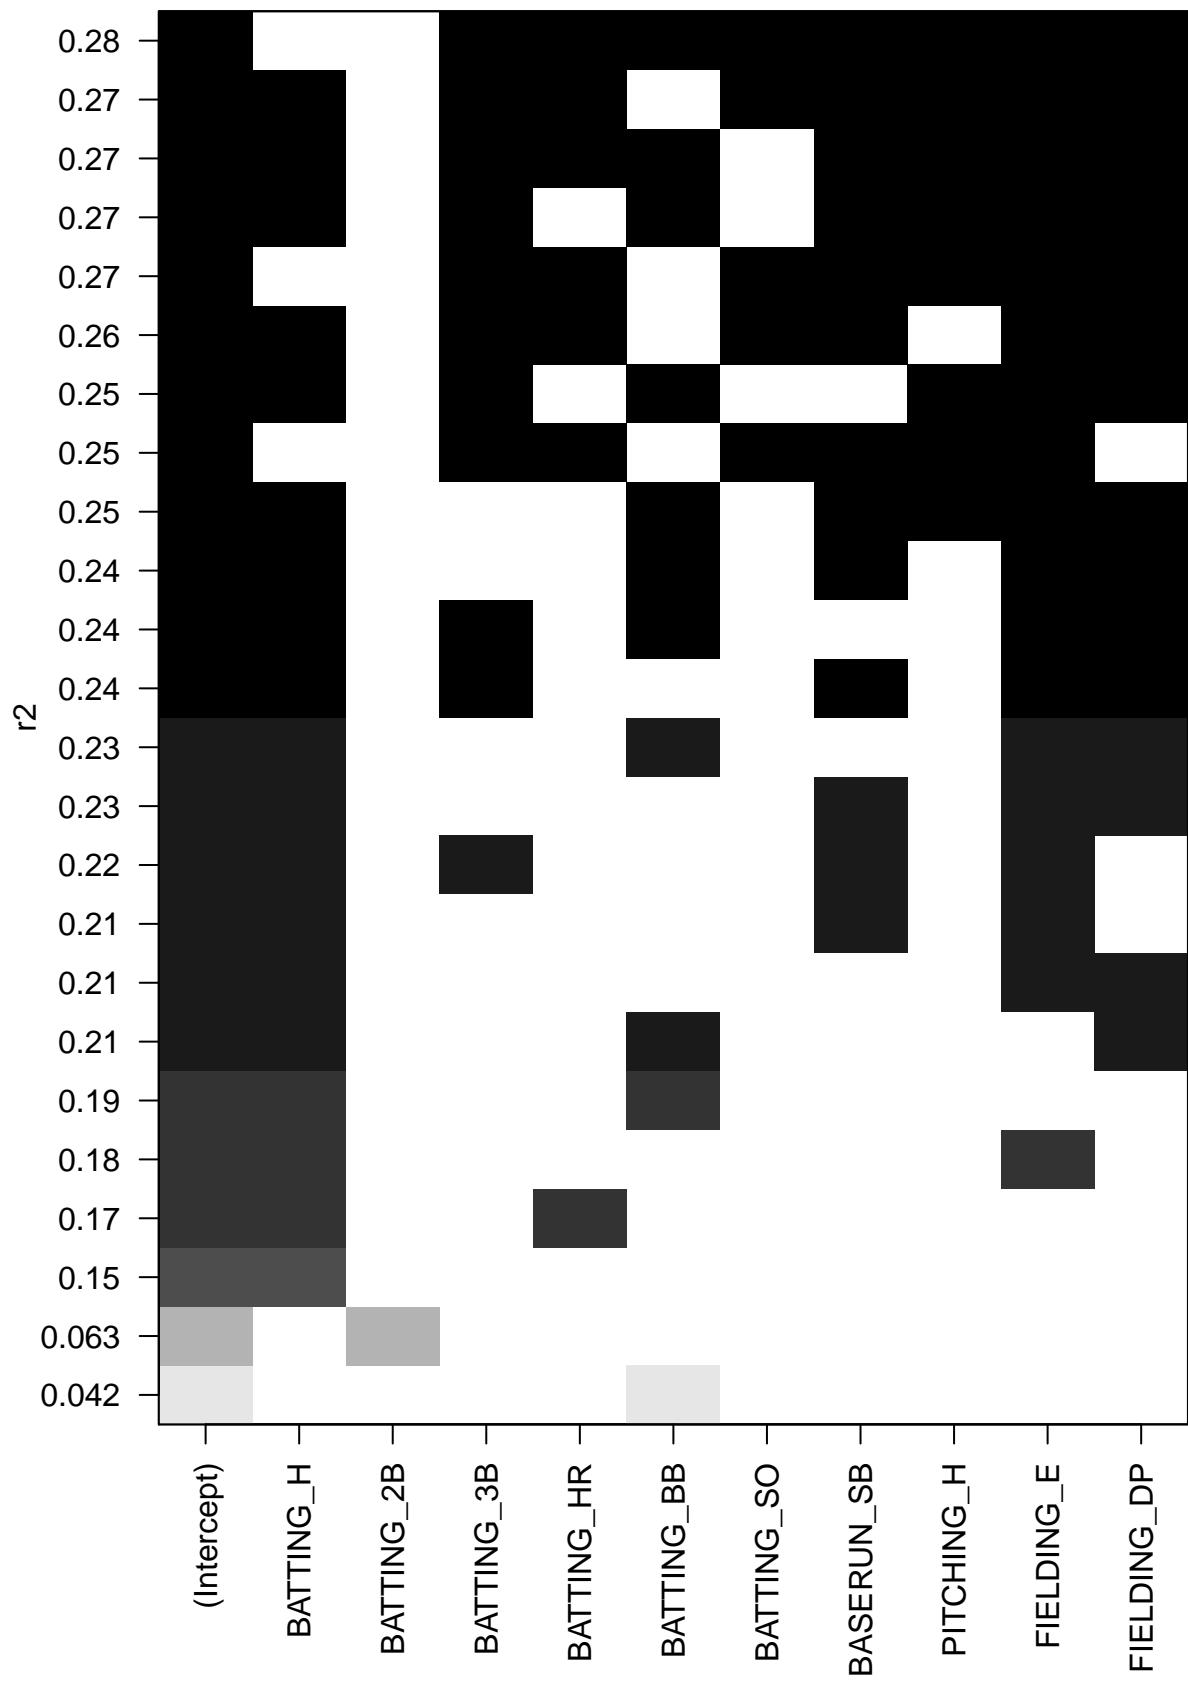
Table 2: Full SCALED Model Coefficients

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	80.971749	0.26662860	304.078090	0.00000000
BATTING_H	3.222999	0.5814092	5.543426	0.00000000
BATTING_2B	-1.364857	0.4226469	-3.229309	0.0012590
BATTING_3B	3.381556	0.4597775	7.354765	0.00000000
BATTING_HR	4.220294	0.6076231	6.945579	0.00000000
BATTING_BB	2.190025	0.3629961	6.033192	0.00000000
BATTING_SO	-2.641712	0.5673232	-4.656450	0.00000034
BASERUN_SB	2.794642	0.3453435	8.092356	0.00000000
PITCHING_H	3.870518	0.5290186	7.316411	0.00000000
FIELDING_E	-7.441986	0.6031058	-12.339437	0.00000000
FIELDING_DP	-2.676186	0.3294610	-8.122924	0.00000000

```
## [1] 0.2888829
```



```
## NULL
```



NULL

3.3.1 Test all of the predictors

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
2229	493421.2	NA	NA	NA	NA
2219	350880.3	10	142541	90.14425	0

3.3.2 Test one predictor

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
2220	355739.4	NA	NA	NA	NA
2219	350880.3	1	4859.127	30.72958	0

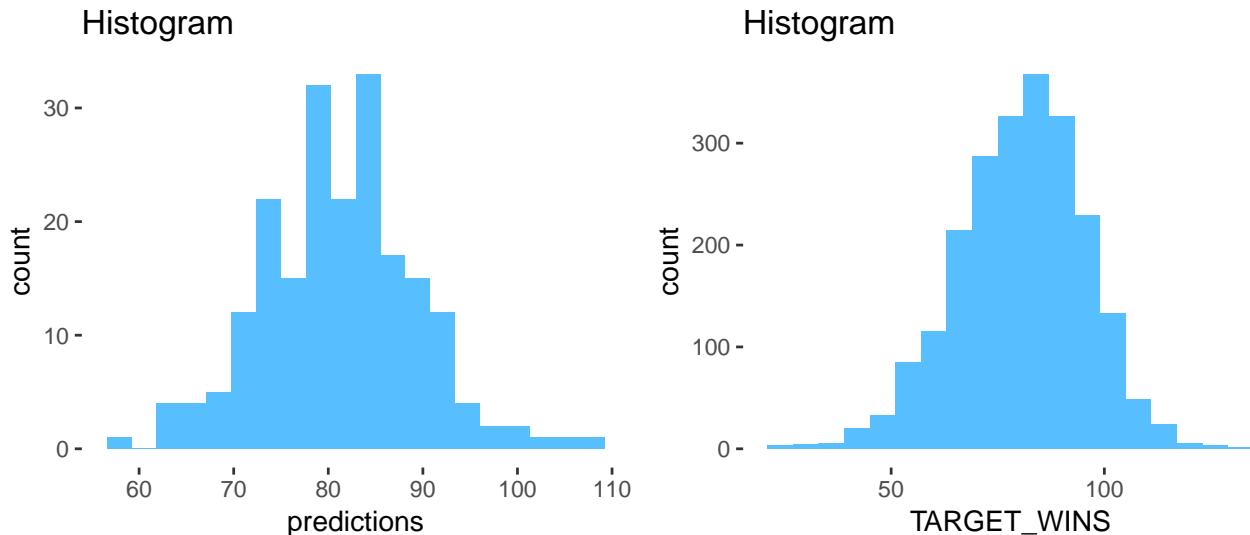
Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
2220	352529.3	NA	NA	NA	NA
2219	350880.3	1	1649.001	10.42844	0.001259

3.3.3 Testing a subspace

```

## 
## Call:
## lm(formula = TARGET_WINS ~ I(BATTING_HR + PITCHING_HR) + I(BATTING_BB +
##     PITCHING_BB) + I(BATTING_SO + PITCHING_SO) + BATTING_H +
##     BATTING_2B + BATTING_3B + BASERUN_SB + BASERUN_CS + PITCHING_H +
##     FIELDING_E + FIELDING_DP, data = all_data)
## 
## Residuals:
##      Min        1Q    Median        3Q       Max
## -67.298   -8.254    0.184    8.251   66.950
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)                33.769875  5.912262  5.712 1.27e-08 ***
## I(BATTING_HR + PITCHING_HR)  0.042078  0.004928  8.538 < 2e-16 ***
## I(BATTING_BB + PITCHING_BB)  0.006255  0.001547  4.044 5.44e-05 ***
## I(BATTING_SO + PITCHING_SO) -0.003880  0.001061 -3.656 0.000262 ***
## BATTING_H                   0.023533  0.004888  4.815 1.57e-06 ***
## BATTING_2B                  -0.007472  0.009482 -0.788 0.430776  
## BATTING_3B                  0.104289  0.017601  5.925 3.61e-09 ***
## BASERUN_SB                  0.023849  0.005918  4.030 5.76e-05 ***
## BASERUN_CS                  0.037198  0.016067  2.315 0.020690 *  
## PITCHING_H                  0.008734  0.001205  7.248 5.80e-13 ***
## FIELDING_E                 -0.035162  0.003271 -10.748 < 2e-16 ***
## FIELDING_DP                 -0.084532  0.013704 -6.168 8.18e-10 ***
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 12.74 on 2218 degrees of freedom
## Multiple R-squared:  0.2704, Adjusted R-squared:  0.2668 
## F-statistic: 74.74 on 11 and 2218 DF,  p-value: < 2.2e-16

```



```

## Start: AIC=11321.3
## TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B + BATTING_HR +
##           BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##           PITCHING_HR + PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP +
##           BP_H + BP_HR + BP_BB + BP_SO
##
##
## Step: AIC=11321.3
## TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B + BATTING_HR +
##           BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##           PITCHING_HR + PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP +
##           BP_H + BP_HR + BP_BB
##
##
## Step: AIC=11321.3
## TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B + BATTING_HR +
##           BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##           PITCHING_HR + PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP +
##           BP_H + BP_HR
##
##
## Step: AIC=11321.3
## TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B + BATTING_HR +
##           BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##           PITCHING_HR + PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP +
##           BP_H
##
##
## Step: AIC=11321.3
## TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B + BATTING_HR +
##           BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##           PITCHING_HR + PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP
##
##             Df Sum of Sq   RSS   AIC
## - PITCHING_HR  1      71.3 352682 11320
## - BATTING_2B   1     147.1 352757 11320
## <none>          352610 11321

```

```

## - BASERUN_CS 1 905.6 353516 11325
## - BATTING_HR 1 932.0 353542 11325
## - BATTING_H 1 1774.0 354384 11330
## - PITCHING_SO 1 1884.8 354495 11331
## - PITCHING_BB 1 2313.2 354924 11334
## - BATTING_SO 1 3164.7 355775 11339
## - BATTING_BB 1 3725.7 356336 11343
## - BASERUN_SB 1 4387.3 356998 11347
## - FIELDING_DP 1 5849.4 358460 11356
## - BATTING_3B 1 7439.0 360049 11366
## - PITCHING_H 1 12791.3 365402 11399
## - FIELDING_E 1 20259.4 372870 11444
##
## Step: AIC=11319.75
## TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B + BATTING_HR +
##           BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##           PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP
##
##             Df Sum of Sq   RSS   AIC
## - BATTING_2B  1    144.5 352826 11319
## <none>          352682 11320
## + PITCHING_HR 1    71.3 352610 11321
## + BP_HR        1    71.3 352610 11321
## - BASERUN_CS  1   923.1 353605 11324
## - BATTING_H   1  1786.1 354468 11329
## - PITCHING_SO 1  1855.1 354537 11329
## - BATTING_SO  1  3113.1 355795 11337
## - PITCHING_BB 1  4312.8 356995 11345
## - BASERUN_SB  1  4384.3 357066 11345
## - FIELDING_DP 1  5907.2 358589 11355
## - BATTING_BB  1  6344.8 359026 11358
## - BATTING_3B  1  7380.9 360063 11364
## - BATTING_HR  1 12111.1 364793 11393
## - PITCHING_H  1 12883.5 365565 11398
## - FIELDING_E  1 21170.8 373853 11448
##
## Step: AIC=11318.66
## TARGET_WINS ~ BATTING_H + BATTING_3B + BATTING_HR + BATTING_BB +
##           BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H + PITCHING_BB +
##           PITCHING_SO + FIELDING_E + FIELDING_DP
##
##             Df Sum of Sq   RSS   AIC
## <none>          352826 11319
## + BATTING_2B  1    144.5 352682 11320
## + PITCHING_HR 1    68.8 352757 11320
## + BP_HR        1    68.8 352757 11320
## - BASERUN_CS  1   997.3 353824 11323
## - PITCHING_SO 1  1748.3 354575 11328
## - BATTING_H   1  1953.4 354780 11329
## - BATTING_SO  1  3092.7 355919 11336
## - PITCHING_BB 1  4257.3 357084 11343
## - BASERUN_SB  1  4451.1 357277 11345
## - FIELDING_DP 1  5838.4 358665 11353
## - BATTING_BB  1  6265.6 359092 11356

```

```

## - BATTING_3B    1    7722.8 360549 11365
## - BATTING_HR   1    12449.8 365276 11394
## - PITCHING_H   1    13383.6 366210 11400
## - FIELDING_E   1    21030.9 373857 11446



| Step          | Df | Deviance  | Resid. Df | Resid. Dev | AIC      |
|---------------|----|-----------|-----------|------------|----------|
|               | NA | NA        | 2215      | 352610.4   | 11321.30 |
| - BP_SO       | 0  | 0.00000   | 2215      | 352610.4   | 11321.30 |
| - BP_BB       | 0  | 0.00000   | 2215      | 352610.4   | 11321.30 |
| - BP_HR       | 0  | 0.00000   | 2215      | 352610.4   | 11321.30 |
| - BP_H        | 0  | 0.00000   | 2215      | 352610.4   | 11321.30 |
| - PITCHING_HR | 1  | 71.32528  | 2216      | 352681.7   | 11319.75 |
| - BATTING_2B  | 1  | 144.53570 | 2217      | 352826.2   | 11318.66 |



##
## Call:
## lm(formula = TARGET_WINS ~ BATTING_H + BATTING_3B + BATTING_HR +
##      BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + PITCHING_H +
##      PITCHING_BB + PITCHING_SO + FIELDING_E + FIELDING_DP, data = imputed_train)
##
## Residuals:
##    Min      1Q  Median      3Q     Max
## -80.887 -8.130   0.041   8.097  68.359
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 33.175729  5.656513  5.865 5.16e-09 ***
## BATTING_H    0.013834  0.003949  3.503 0.000468 ***
## BATTING_3B   0.125448  0.018008  6.966 4.28e-12 ***
## BATTING_HR   0.090183  0.010196  8.845 < 2e-16 ***
## BATTING_BB   0.083849  0.013363  6.275 4.20e-10 ***
## BATTING_SO  -0.029123  0.006606 -4.408 1.09e-05 ***
## BASERUN_SB   0.032719  0.006187  5.289 1.35e-07 ***
## BASERUN_CS   0.040007  0.015982  2.503 0.012375 *
## PITCHING_H   0.014048  0.001532  9.170 < 2e-16 ***
## PITCHING_BB -0.060007  0.011602 -5.172 2.52e-07 ***
## PITCHING_SO  0.017834  0.005381  3.314 0.000933 ***
## FIELDING_E   -0.037924  0.003299 -11.496 < 2e-16 ***
## FIELDING_DP -0.082243  0.013579 -6.057 1.63e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.62 on 2217 degrees of freedom
## Multiple R-squared:  0.2849, Adjusted R-squared:  0.2811
## F-statistic: 73.62 on 12 and 2217 DF, p-value: < 2.2e-16
##
## Call:
## lm(formula = TARGET_WINS ~ BATTING_3B + BATTING_HR + BATTING_BB +
##      PITCHING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS + log(PITCHING_H) +
##      log(FIELDING_E) + FIELDING_DP, data = imputed_train)
##
## Residuals:
##    Min      1Q  Median      3Q     Max
## -65.732 -7.705   0.020   7.770  70.422

```

```

## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|) 
## (Intercept)           -1.565e+02  2.608e+01 -6.001 2.28e-09 ***
## BATTING_3B            1.598e-01  1.635e-02  9.771 < 2e-16 ***
## BATTING_HR             6.900e-02  9.843e-03  7.010 3.14e-12 ***
## BATTING_BB             7.055e-02  1.446e-02  4.880 1.14e-06 ***
## PITCHING_BB           -7.169e-02  9.991e-03 -7.175 9.79e-13 ***
## BATTING_SO             -1.133e-02  2.130e-03 -5.319 1.15e-07 ***
## BASERUN_SB              2.978e-02  5.672e-03  5.251 1.66e-07 ***
## BASERUN_CS              5.396e-02  1.574e-02  3.427 0.000621 *** 
## log(PITCHING_H)        4.287e+01  3.172e+00 13.517 < 2e-16 ***
## log(FIELDING_E)        -1.604e+01  1.062e+00 -15.101 < 2e-16 ***
## FIELDING_DP             -9.269e-02  1.322e-02 -7.012 3.11e-12 ***
## BATTING_BB:PITCHING_BB  2.575e-05  1.544e-05  1.668 0.095489 . 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 12.4 on 2218 degrees of freedom
## Multiple R-squared:  0.3084, Adjusted R-squared:  0.3049 
## F-statistic:  89.9 on 11 and 2218 DF,  p-value: < 2.2e-16

```

3.4 MODEL 3

Jeremy: Created base case for comparison.

```

## 
## Call:
## lm(formula = TARGET_WINS ~ +BATTING_3B + BATTING_HR + BATTING_BB +
##      BATTING_SO + BASERUN_SB + PITCHING_H + PITCHING_BB + PITCHING_SO +
##      FIELDING_E + FIELDING_DP, data = imputed_train)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -92.817  -8.348   0.121   8.022  73.344 
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|) 
## (Intercept)          50.683827  3.800008 13.338 < 2e-16 ***
## BATTING_3B            0.155798  0.016019  9.726 < 2e-16 ***
## BATTING_HR             0.096477  0.008908 10.830 < 2e-16 ***
## BATTING_BB             0.093546  0.013124  7.128 1.37e-12 ***
## BATTING_SO             -0.032828  0.006567 -4.999 6.23e-07 ***
## BASERUN_SB              0.045900  0.004783  9.596 < 2e-16 ***
## PITCHING_H             0.016998  0.001182 14.377 < 2e-16 ***
## PITCHING_BB            -0.069058  0.011406 -6.054 1.65e-09 ***
## PITCHING_SO             0.018501  0.005379  3.440 0.000593 *** 
## FIELDING_E             -0.040897  0.003115 -13.130 < 2e-16 ***
## FIELDING_DP            -0.084052  0.012936 -6.498 1.00e-10 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 12.66 on 2219 degrees of freedom
## Multiple R-squared:  0.2793, Adjusted R-squared:  0.2761

```

```
## F-statistic: 86 on 10 and 2219 DF, p-value: < 2.2e-16
```

Jeremy:

- The p-value of our F-stat is below .05 so we reject null that coefficients of zero better fit.
 - AR² indicates the model is explaining about 30% of the variation in TARGET_WINS
 - Based on high p-vals, removed variables BATTING_H, BATTING_2B, BASERUN_CS, PITCHING_HR through backward elimination.
 - Negative coefficients of BATTING_SO, PITCHING_BB, FIELDING_E, FIELDING_DP accord with intuition that would drag rather than power wins.
 - 3Bs and balls do more to drive wins than homeruns and strikeouts.
-

Jeremy: since the off-def transform didn't yield improvements on base case, explored log transform model for right-skewed independent variables - PITCHING_H, PITCHING_HR, PITCHING_SO.

```
##  
## Call:  
## lm(formula = TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B +  
##      BATTING_HR + BATTING_BB + BATTING_SO + BASERUN_SB + log(PITCHING_H) +  
##      PITCHING_BB + FIELDING_E + FIELDING_DP, data = imputed_train)  
##  
## Residuals:  
##      Min        1Q    Median        3Q       Max  
## -78.012   -8.149    0.253    7.985   73.220  
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) -4.050e+02  3.700e+01 -10.947 < 2e-16 ***  
## BATTING_H   -1.346e-02  6.034e-03 -2.232  0.02574 *  
## BATTING_2B  -1.500e-02  9.146e-03 -1.641  0.10102  
## BATTING_3B   1.405e-01  1.766e-02  7.960 2.72e-15 ***  
## BATTING_HR   7.230e-02  9.704e-03  7.451 1.32e-13 ***  
## BATTING_BB   1.247e-01  1.266e-02  9.846 < 2e-16 ***  
## BATTING_SO  -6.554e-03  2.330e-03 -2.813  0.00495 **  
## BASERUN_SB   4.380e-02  4.763e-03  9.197 < 2e-16 ***  
## log(PITCHING_H) 6.866e+01  5.760e+00 11.921 < 2e-16 ***  
## PITCHING_BB -9.518e-02  1.063e-02 -8.953 < 2e-16 ***  
## FIELDING_E  -4.740e-02  3.550e-03 -13.350 < 2e-16 ***  
## FIELDING_DP -8.853e-02  1.294e-02 -6.842 1.01e-11 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 12.49 on 2218 degrees of freedom  
## Multiple R-squared:  0.2991, Adjusted R-squared:  0.2956  
## F-statistic: 86.03 on 11 and 2218 DF, p-value: < 2.2e-16
```

Jeremy - F-stat is significant but lower than base case model; however, AR² has risen marginally with inclusion of the log-transformed pitching variables. - Interestingly, BATTING_H and BATTING_2B p-vals suggest including in this model (removed from base case); however, based on their coefficients hits and second base runs negatively impact wins, which seems counterintuitive. - Based on high p-vals, removed BASERUN_CS, PITCHING_HR, PITCHING_SO; the log transforms of those two pitching variables did not bring p-val below .05. - Crazy different intercept - plot to visualize.

Jeremy: Created blend of two different transforms - the offdef and log. Several iterations have not yielded

better F-stats or R², but need to investigate further - especially why certain variables can sit alongside originals when this caused conflict before.

```

## 
## Call:
## lm(formula = TARGET_WINS ~ BATTING_H + BATTING_2B + BATTING_3B +
##      BATTING_HR + BATTING_BB + BATTING_SO + BASERUN_SB + BASERUN_CS +
##      PITCHING_H + log(PITCHING_HR + 1e-04) + PITCHING_BB + log(PITCHING_SO +
##      1e-04) + FIELDING_E + FIELDING_DP + BP_H, data = imputed_train)
## 
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -77.344 -8.117 -0.027  7.965 69.900 
## 
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 35.568590  6.460266  5.506 4.10e-08 ***
## BATTING_H    0.015810  0.005011  3.155 0.00163 ** 
## BATTING_2B   -0.004716  0.009341 -0.505 0.61368  
## BATTING_3B    0.132683  0.018302  7.249 5.75e-13 ***
## BATTING_HR   0.084228  0.011471  7.343 2.93e-13 *** 
## BATTING_BB   0.044636  0.009638  4.631 3.85e-06 *** 
## BATTING_SO   -0.006234  0.002487 -2.507 0.01226 *  
## BASERUN_SB   0.025473  0.005999  4.246 2.26e-05 *** 
## BASERUN_CS   0.044209  0.016681  2.650 0.00810 ** 
## PITCHING_H   0.012942  0.001608  8.049 1.35e-15 *** 
## log(PITCHING_HR + 1e-04) 0.281261  0.560322  0.502 0.61574  
## PITCHING_BB  -0.023742  0.008032 -2.956 0.00315 ** 
## log(PITCHING_SO + 1e-04) -0.802335  0.335357 -2.392 0.01682 *  
## FIELDING_E   -0.036429  0.003252 -11.201 < 2e-16 *** 
## FIELDING_DP  -0.085646  0.013724 -6.241 5.21e-10 *** 
## BP_H          NA        NA        NA        NA      
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 12.64 on 2215 degrees of freedom
## Multiple R-squared:  0.2833, Adjusted R-squared:  0.2788 
## F-statistic: 62.55 on 14 and 2215 DF,  p-value: < 2.2e-16

```

4 SELECT MODELS

4.1 Instructions:

Decide on the criteria for selecting the best multiple linear regression model. Will you select a model with slightly worse performance if it makes more sense or is more parsimonious? Discuss why you selected your model. For the multiple linear regression model, will you use a metric such as Adjusted R², RMSE, etc.? Be sure to explain how you can make inferences from the model, discuss multi-collinearity issues (if any), and discuss other relevant model output. Using the training data set, evaluate the multiple linear regression model based on (a) mean squared error, (b) R², (c) F-statistic, and (d) residual plots. Make predictions using the evaluation data set.